

Deadlock (1)

Dave Eckhardt
Garth Gibson
Bruce Maggs
Geoff Langdale

Synchronization – P2

- You should *really* have, today:
 - Figured out where wrappers belong, why
 - Made some system calls
 - Designed mutexes & condition variables
 - Drawn pictures of thread stacks (even if not perfect)
- Wednesday:
 - Coded mutexes and condition variables
 - Thoughtful design for `thr_create()`, maybe `thr_join()`
 - Some code for `thr_create()`, and some “experience”
 - The `startle` test running

Debugging Reminder

- We can't really help with queries like:
 - We did x... then something strange happened...
 - ...can you tell us why?
- You need to progress beyond “*something* happened”
 - What happened, exactly?
 - printf() is not always the right tool
 - output correct only if run-time environment is right
 - captures only what you told it to, only “C-level” stuff
 - *changes your code* by its mere presence!!!
 - We're serious about examining register dumps!
 - Overall, maybe re-read “Debugging” lecture notes

Synchronization – Readings

- **Next three lectures**
 - **Deadlock: 6.5.3, 6.6.3, Chapter 7**
- **Reading ahead**
 - **Scheduling: Chapter 5**
 - **Virtual Memory: Chapter 8, Chapter 9**
- **Don't forget about reading list on web site**

Outline

- Process resource graph
- What is deadlock?
- Deadlock *prevention*
- Next time
 - Deadlock *avoidance*
 - Deadlock *recovery*

Tape Drives

- A word on “tape drives”
 - Ancient computer resources
 - Access is sequential, read/write
 - Any tape can be mounted on any drive
 - One tape at a time is mounted on a drive
 - Doesn't make sense for multiple processes to simultaneously access a drive
 - Reading/writing a tape takes a while
- Think “CD burner” ...




IBM 3420 (1970-1987)
www.ibm.com/ibm/history
Not for publication




Data General 6023
wps.com/NOVA4


Process/Resource graph




Process/Resource graph




Waiting




Release



Reallocation



Multi-instance Resources



Definition of Deadlock

- A deadlock
 - Set of N processes
 - Each waiting for an event
 - ...which can be caused *only by another process in the set*
- Every process will wait forever

Deadlock Examples

- **Simplest form**
 - Process 1 owns printer, wants tape drive
 - Process 2 owns tape drive, wants printer
- **Less-obvious**
 - Three tape drives
 - Three processes
 - Each has one tape drive
 - Each wants “just” one more
 - Can't blame anybody, but problem is still there

Deadlock Requirements

- **Mutual Exclusion**
- **Hold & Wait**
- **No Preemption**
- **Circular Wait**

Mutual Exclusion

- Resources aren't “thread-safe” (“reentrant”)
- Must be allocated to one process/thread at a time
- Can't be shared
 - Programmable Interrupt Timer
 - Can't have a different reload value for each process

Hold & Wait

- Process holds some resources while waiting for more

```
mutex_lock( &m1 );
```

```
mutex_lock( &m2 );
```

```
mutex_lock( &m3 );
```

- This locking behavior is *typical*


No Preemption

- Can't force a process to give up a resource
- Interrupting a CD-R burn creates a “coaster”
 - So don't do that
- Obvious solution
 - CD-R device driver forbids second simultaneous `open()`
 - If you can't open it, you can't pre-empt it...

Circular Wait

- **Process 0 needs something process 4 has**
 - **Process 4 needs something process 7 has**
 - **Process 7 needs something process 1 has**
 - **Process 1 needs something process 0 has – uh-oh...**
- **Described as “cycle in the resource graph”**


Cycle in Resource Graph




Deadlock Requirements

- Mutual Exclusion
- Hold & Wait
- No Preemption
- Circular Wait
- *Each deadlock* requires *all four*


Multi-Instance Cycle



Multi-Instance Cycle *(With Rescuer!)*




Cycle Broken



Dining Philosophers

- **The scene**
 - 410 staff members at a Chinese restaurant
 - A little short on utensils

Dining Philosophers



Dining Philosophers

- **Processes**
 - 5, one per person
- **Resources**
 - 5 bowls (dedicated to a diner: no contention: ignore)
- **5 chopsticks**
 - 1 between every adjacent pair of diners
- **Contrived example?**
 - Illustrates contention, starvation, deadlock


Dining Philosophers

- **A simple rule for eating**
 - **Wait until the chopstick to your right is free; take it**
 - **Wait until the chopstick to your left is free; take it**
 - **Eat for a while**
 - **Put chopsticks back down**


Dining Philosophers Deadlock

- **Everybody reaches right...**
 - ...at the same time?


Reaching Right



Successful Acquisition



Deadlock!



Dining Philosophers – State

```
int stick[5] = { -1 }; /* owner */
condition avail[5]; /* newly avail. */
mutex table = { available };

/* Right-handed convention */
right = diner;          /* 3 ⇒ 3 */
left = (diner + 4) % 5; /* 3 ⇒ 7 ⇒ 2 */
```

start_eating(int diner)

```
mutex_lock(table);
```

```
while (stick[right] != -1)  
    condition_wait(avail[right], table);  
stick[right] = diner;
```

```
while (stick[left] != -1)  
    condition_wait(avail[left], table);  
stick[left] = diner;
```

```
mutex_unlock(table);
```

done_eating(int diner)

```
mutex_lock(table);
```

```
stick[left] = stick[right] = -1;
```

```
condition_signal(avail[right]);
```

```
condition_signal(avail[left]);
```

```
mutex_unlock(table);
```


Can We Deadlock?

- **At first glance the table mutex protects us**
 - **Can't have “everybody reaching right at same time”...**
 - **...mutex means only one person can access table...**
 - **...so allows only one reach at the same time, right?**


Can We Deadlock?

- At first glance the table mutex protects us
 - Can't have “everybody reaching right at same time”...
 - ...mutex means only one person can access table...
 - ...so allows only one reach at the same time, right?
- Maybe we can!
 - *condition_wait()* is a “reach”
 - Can everybody end up in *condition_wait()*?


First diner gets both chopsticks




Next gets right, waits on left




Next two get right, wait on left




Last waits on right




First diner stops eating - *briefly*



First diner stops eating - *briefly*




Next Step – *One* Possibility




“Natural” –
longest-waiting diner progresses

Next Step – *Another* Possibility




Or –
somebody else!


Last diner gets right, waits on left



First diner gets right, waits on left



Now things get boring



Deadlock - What to do?

- **Prevention**
- **Avoidance**
- **Detection/Recovery**
- **Just reboot when it gets “too quiet”**

1: Prevention

- **Restrict behavior or resources**
 - Find a way to violate one of the 4 conditions
 - To wit...?
- **What we will talk about today**
 - 4 conditions, 4 possible ways

2: Avoidance

- Processes *pre-declare* usage patterns
- Dynamically examine requests
 - Imagine what other processes could ask for
 - Keep system in “safe state”

3: Detection/Recovery

- Maybe deadlock won't happen today...
- ...Hmm, it seems quiet...
- ...Oops, here is a cycle...
- *Abort some process*
 - Ouch!

4: Reboot When It Gets “Too Quiet”

- **Which systems would be so simplistic?**

Four Ways to Forgiveness

- *Each deadlock* requires *all four*
 - Mutual Exclusion
 - Hold & Wait
 - No Preemption
 - Circular Wait
- “Deadlock Prevention” - this is a technical term
 - *Pass a law* against one (pick one)
 - Deadlock happens only if somebody *transgresses!*

Outlaw Mutual Exclusion?

- Approach: *ban* single-user resources
 - Require all resources to “work in shared mode”
- Problem
 - Chopsticks???
 - Many resources don't work that way

Outlaw Hold&Wait?

- Acquire resources *all-or-none*

```
start_eating(int diner)
```

```
mutex_lock(table);
```

```
while (1)
```

```
    if (stick[lr] == stick[rt] == -1)
```

```
        stick[lr] = stick[rt] = diner
```

```
        mutex_unlock(table)
```

```
        return;
```

```
    condition_wait(released, table);
```


Problems

- **“Starvation”**
 - Larger resource set makes grabbing everything harder
 - No guarantee a diner eats in bounded time
- **Low utilization**
 - Larger peak resource needs hurts whole system always
 - Must allocate 2 chopsticks (and waiter!)
 - Nobody else can use waiter while you eat

Outlaw Non-preemption?

- Steal resources from sleeping processes!

```
start_eating(int diner)
right = diner;    rright = (diner+1)%5;
mutex_lock(table);
while (1)
    if (stick[right] == -1)
        stick[right] = diner
    else if (stick[rright] != rright)
        /* right person can't be eating: take! */
        stick[right] = diner;
...same for left...wait() if must...
mutex_unlock(table);
```

Problem

- **Some resources cannot be cleanly preempted**
 - CD burner

Outlaw Circular Wait?

- Impose *total order* on all resources
- Require acquisition in *strictly increasing order*
 - Static order may work: allocate memory, then files
 - Dynamic – may need to “start over” sometimes
 - Traversing a graph
 - lock(4), visit(4) /* 4 has an edge to 13 */
 - lock(13), visit(13) /* 13 has an edge to 0 */
 - lock(0)?
 - Nope!
 - unlock(4), unlock(13)
 - lock(0), lock(4), lock(13), ...

Assigning Diners a Total Order

- Lock order: 4, 3, 2, 1, 0 \equiv right chopstick, then left
 - Diner 4 \Rightarrow lock(4); lock(3);
 - Diner 3 \Rightarrow lock(3); lock(2);

Assigning Diners a Total Order

- Lock order: 4, 3, 2, 1, 0 \equiv right chopstick, then left
 - Diner 4 \Rightarrow lock(4); lock(3);
 - Diner 3 \Rightarrow lock(3); lock(2);
 - Diner 0 \Rightarrow lock(0); lock(4); /* invalid lock order! */
 - Requires special-case locking code to get order right

```
if diner == 0
    right = (diner + 4) % 5;
    left = diner;
else
    right = diner;
    left = (diner + 4) % 5;
...
```

Problem

- **May not be possible to force allocation order**
 - **Some trains go east, some go west**

Deadlock Prevention problems

- Typical resources *require* mutual exclusion
- All-at-once allocation can be *painful*
 - Hurts efficiency
 - May starve
 - Resource needs may be unpredictable
- Preemption may be *impossible*
 - Or may lead to starvation
- Ordering restrictions may be *impractical*

Deadlock Prevention

- Pass a law against one of the four ingredients
 - Great if you can find a tolerable approach
- **Very** tempting to just let processes try their luck

Deadlock is not...

- **...a simple synchronization bug**
 - Deadlock remains even when those are cleaned up
 - Deadlock is a resource usage design problem
- **...the same as starvation**
 - Deadlocked processes don't ever get resources
 - Starved processes don't ever get resources
 - Deadlock is a “progress” problem; starvation is a “bounded waiting” problem
- **....that “after-you, sir” dance in the corridor**
 - That's “livelock” – continuous changes of state without forward progress

Next Time

- **Deadlock Avoidance**
- **Deadlock Recovery**