

**15-410**

*“...Does this look familiar?...”*

# **File System (Internals)**

**Mar. 23, 2011**

**Dave Eckhardt**

**Garth Gibson**

**Greg Hartman**

# Synchronization

## Today

- Chapter 11 (not: Log-structured, NFS, WAFL)

# Outline

**File system code layers (abstract)**

**Disk, memory structures**

**Unix “VFS” layering indirection**

**Directories**

**Block allocation strategies, free space**

**Cache tricks**

**Recovery, backups**

# File System Layers

## Device drivers

- `read/write(disk, start-sector, count)`

## Block I/O

- `read/write(partition, block) [cached]`

## File I/O

- `read/write(file, block)`

## File system

- manage directories, free space

# File System Layers

## Multi-filesystem namespace

- Partitioning, names for devices
- Mounting
- Unifying multiple file system *types*
  - UFS, ext2fs, ext3fs, zfs, FAT, 9660, ...

# Shredding Disks

## Split disk into *partitions*/slices/minidisks/...

- MBR (PC): 4 “partitions” – Windows, FreeBSD, Plan 9, ...
- APM (Mac): “volumes” – can split: OS 9, OS X, user files
- GPT (new, multi-platform) - many partitions, long names

## Or: glue disks together into *volumes*/logical disks

## A partition (of a disk or of a volume) may contain...

- Paging area
  - Indexed by in-memory structures
  - “random garbage” when OS shuts down
- File system
  - Block allocation: file #  $\Rightarrow$  block list
  - Directory: name  $\Rightarrow$  file #

# Shredding Disks

```
# fdisk -s
```

```
/dev/ad0: 993 cyl 128 hd 63 sec
```

Part	Start	Size	Type	Flags
1:	63	1233729	0x06	0x00
2:	1233792	6773760	0xa5	0x80

(A 4-gigabyte disk)

# Shredding Disks

8 partitions:

#	size	offset	fstype	[fsize	bsize	bps/cpg]		
a:	131072	0	4.2BSD	2048	16384	101	# (Cyl.	0 - 16*)
b:	393216	131072	swap				# (Cyl.	16*- 65*)
c:	6773760	0	unused	0	0		# (Cyl.	0 - 839)
e:	65536	524288	4.2BSD	2048	16384	104	# (Cyl.	65*- 73*)
f:	6183936	589824	4.2BSD	2048	16384	89	# (Cyl.	73*- 839*)

Filesystem	1K-blocks	Used	Avail	Capacity	Mounted on
/dev/ad0s2a	64462	55928	3378	94%	/
/dev/ad0s2f	3043806	2608458	191844	93%	/usr
/dev/ad0s2e	32206	7496	22134	25%	/var
procfs	4	4	0	100%	/proc

(FreeBSD 4.7 on ThinkPad 560X)



# Disk Structures

## **Boot area (first block/track/cylinder)**

- Interpreted by hardware bootstrap (“BIOS”)
- May include partition table

## **File system control block**

- Key parameters: #blocks, metadata layout
- Unix: “superblock”

## **“File control block” (Unix: “inode”)**

- ownership/permissions
- data location

## **Possibly a free-space map as well**

# Memory Structures

## **In-memory partition tables**

- Sanity check file system I/O fits in correct partition

## **Cached directory information**

## **System-wide open-file table**

- In-memory file control blocks

## **Process open-file tables**

- Open mode (read/write/append/...)
- “Cursor” (read/write position)

# VFS layer

## Goals

- Allow one machine to use multiple file system *types*
  - Unix FFS
  - MS-DOS FAT
  - CD-ROM ISO9660
  - Remote/distributed: NFS/AFS
- Standard system calls should work transparently

## Solution?

# VFS layer

## Goals

- Allow one machine to use multiple file system *types*
  - Unix FFS
  - MS-DOS FAT
  - CD-ROM ISO9660
  - Remote/distributed: NFS/AFS
- Standard system calls should work transparently

## Solution

- Insert a level of indirection!

# Single File System

```
n = read(fd, buf, size)
```

```
INT 54
```

```
sys_read(fd, buf, len)
```

```
namei()
```

```
iget()
```

```
iput()
```

```
sleep()
```

```
rdblkl(dev, N)
```

```
wakeup()
```

```
startIDE()
```

```
IDEintr()
```

# VFS “Virtualization”

```
n = read(fd, buf, size)
```

```
INT 54
```

```
namei()
```

```
vfs_read()
```

```
ufs_read()
```

```
procfs_read()
```

```
ufs_lookup()
```

```
procfs_domem()
```

```
iget()
```

```
iput()
```

# VFS layer – file system operations

**These operate on file *systems*, not individual files**

```
struct vfsops {  
    char *name;  
    int (*vfs_mount)();  
    int (*vfs_statfs)();  
    int (*vfs_vget)();  
    int (*vfs_unmount)();  
    ...  
}
```

# VFS layer – file operations

## **Each VFS provides an array of per-file methods**

- `VOP_LOOKUP(vnode, new_vnode, name)`
- `VOP_CREATE(vnode, new_vnode, name, attributes)`
- `VOP_OPEN(vnode, mode, credentials, process)`
- `VOP_READ(vnode, uio, readwrite, credentials)`

## **Operating system provides fs-independent code**

- Validating system call parameters
- Moving data from/to user memory
- Thread sleep/wakeup
- Caches (data blocks, name  $\Rightarrow$  vnode mappings)



# Directories

**Old: one namei()  $\Rightarrow$  VFS: fs-provided vnode method**

- `vnode2 = VOP_LOOKUP(vnode1, name)`

**Traditional Unix FFS directories**

- List of (name,inode #) - not sorted!
- Names are variable-length
- Lookup is linear
  - How long does it take to delete N files?

**Common alternative: hash-table directories**

# Allocation / Mapping

## Allocation problem

- Where do I put the next block of this file?
  - “Near the previous block” is not a bad idea
  - Beyond that, it gets complicated

## Mapping problem

- Where was block 32 of this file previously put?
- Similar to virtual memory
  - Multiple large “address spaces” *specific to each file*
  - Only one underlying “address space” of blocks
  - Source address space may be sparse!

# Allocation / Mapping

**Contiguous**

**Linked**

**FAT**

**Indexed**

Linked

Multi-level

Unix (index tree)

# Allocation – Contiguous

## Approach

- File location defined as (start, length)

## Motivation

- Sequential disk accesses are cheap
- Bookkeeping is easy

## Issues

- Dynamic storage allocation (fragmentation, compaction)
- Must pre-declare file size at creation
- This should sound familiar

# Allocation – Linked

## Approach

- File location defined as (start)
- Each disk block contains pointer to next block

## Motivation

- Avoids fragmentation problems
- Allows file growth

## Issues?

# Allocation – Linked

## Issues

- 508-byte blocks don't match memory pages
- In general, one seek per block read/written - *slow!*
- *Very* hard to access file blocks at random
  - `lseek(fd, 37 * 1024, SEEK_SET);`

## Benefit

- Can recover files even if directories destroyed

## Common modification

- Link multi-block *clusters*, not blocks

# Allocation – FAT

## **Used by MS-DOS, OS/2, Windows**

- Digital cameras, GPS receivers, printers, PalmOS, ...

## **Semantically same as linked allocation**

But next-block links stored “out of band” in a table

- Result: nice 512-byte sectors for data

## **Table at start of disk**

- Next-block pointer array
- Indexed by block number
- Next=0 means “free”

# Allocation – FAT

7
2
5
-1
3
-1
0
-1

hello.jav	0
dir. c	1
sys.ini	4




# Allocation - FAT


7
2
5
-1
3
-1
0
-1

hello.jav	0
dir. c	1
sys.ini	4

# Allocation - FAT



# Allocation - FAT



# Allocation – FAT

## Issues

- Damage to FAT scrambles entire file system
  - Solution: mirror the FAT
- Generally *two* seeks per block read/write
  - Seek to FAT, read, seek to actual block (repeat)
  - Unless FAT can be cached well in RAM
- Still *somewhat* hard to access random file blocks
  - Linear time to walk through FAT
- FAT may be a “hot spot” (everybody needs to access it)
- Lots of FAT updates (near beginning of disk)
  - Even if files being modified are far away

# Allocation – Indexed

## Motivation

- Avoid fragmentation problems
- Allow file growth
- *Improve random access*

## Approach

- *Per-file* block array
- File block number indexes into table, yields disk block number
- No  $O(n)$  sequential steps

99	3004
100	-1
101	-1
3001	-1
3002	6002
-1	-1
-1	-1
-1	-1

# Allocation – Indexed

## Allows “holes”

- `foo.c` is sequential
- `foo.db`, blocks 1..3  $\Rightarrow$  -1
  - logically “blank”

## “sparse allocation”

- a.k.a. “holes”
- `read()` returns nulls
- `write()` requires alloc
- file “size”  $\neq$  file “size”
  - `ls -l` index of last byte
  - `ls -ls` number of blocks

foo.c	foo.db
99	3004
100	-1
101	-1
3001	-1
3002	6002
-1	-1
-1	-1
-1	-1

# Allocation – Indexed

## How big should index block be?

- Too small: limits file size
- Too big: lots of wasted pointers

## Combining index blocks

- Linked
- Multi-level
- What Unix actually does


# Linked Index Blocks

**Last pointer indicates next index block**

**Simple**

**Access is not-so-random**

- $O(n/c)$  is still  $O(n)$
- $O(n)$  *disk transfers*






# Multi-Level Index Blocks

**Index blocks of index  
blocks**

**Does this look familiar?**

**Allows *big* holes**



# Unix Index Blocks


## Intuition

- *Many* files are small
  - Length = 0, length = 1, length < 80, ...
- Some files are *huge* (gigabytes... maybe terabytes!)


## “Clever heuristic” in Unix FFS inode

- inode struct contains 12 “direct” block pointers
  - 12 block numbers \* 8 KB/block = 96 KB
  - Availability is “free” - must read inode to open() file anyway
- inode struct also contains 3 indirect block pointers
  - single-indirect, double-indirect, triple-indirect


# Unix Index Blocks




# Unix Index Blocks




# Unix Index Blocks



# Unix Index Blocks



# Unix Index Blocks



Triple indirect can address  $\gg 2^{32}$  bytes

# Tracking Free Space

## Bit-vector

- 1 bit per block: boolean “free”
- Check each word vs. 0
- Use “first bit set” instruction
- Text example
  - 1.3 GB disk, 512 B sectors: 332 KB bit vector

**Need to keep (much of) it in RAM**



# Tracking Free Space

## Linked list?

- Superblock points to first free block
- Each free block points to next

## Cost to allocate N blocks is linear

- Free block can point to *multiple* free blocks
  - 512 bytes = 128 (4-byte) block numbers
- FAT approach provides free-block list “for free”

## Keep free-*extent* lists

- (block, sequential-block-count)

# Unified Buffer Cache

## Traditional two-cache approach

- Page cache, file-system cache often totally independent
  - Page cache chunks according to hardware page size
  - File cache chunks according to “file system block” size
  - Different code, different RAM pools
- How much RAM to devote to each one?

## Observation

- Why not have just one cache?
  - Mix automatically varies according to load
    - » “cc” wants more disk cache
    - » Firefox wants more VM cache

# Unified Buffer Cache - Warning!

## **“Virtual memory architecture in SunOS”**

Gingell, Moran, & Shannon

*USENIX 1987 Summer Conference*

“The work has consumed approximately four man-years of effort over a year and a half of real time. A surprisingly large amount of effort has been drained by efforts to interpose the VM system as the logical cache manager for the file systems...”

# Cache tricks

## Read-ahead

```
for (i = 0; i < filesize; ++i)
    putc(getc(infile), outfile);
```

- System observes sequential reads
  - File block 0, 1, 2, ...
  - Can pipeline reads to overlap “computation”, read latency
    - » Request for block 2 triggers disk read of block 3

## Free-behind / replace-behind

- Discard buffer from cache when next is requested
- Good for large files
- “Anti-LRU” (evict “MRU” instead of “LRU”)

# Recovery

## System crash...now what?

- Some RAM contents were lost
- Free-space list on disk may be wrong
- Scan file system
  - Check invariants
    - » Unreferenced files
    - » Double-allocated blocks
    - » Unallocated blocks
  - Fix problems
    - » Expert user???

## Modern approach

- “Journal” changes (see upcoming Transactions lecture)

# Backups

## **Incremental “Towers of Hanoi” approach - traditional**

- **Monthly: dump entire file system**
- **Weekly: dump changes since last monthly**
- **Daily: dump changes since last weekly**
- **Restore a file?**
  - **Most-recent “monthly” tape definitely has a copy**
  - **Any one of the “weekly” tapes might have a copy (scan all)**
  - **Any one of the “daily” tapes might have a copy (scan all)**

# Backups

## **Merge approach (“TiBS”) - [www.teradactyl.com](http://www.teradactyl.com)**

- **Collect changes since yesterday**
  - **Scan file system by modification time**
- **Two tape drives merge yesterday's tape, today's delta, yielding a new tape**
  - **Tape drives write much faster when they're “streaming” (vs. start/stop), so “extra work” may not hurt much**
- **Keep as many tapes as you want to, recycle the rest**
- **Restoring is fast (stream one tape onto disks)**
- **Files stored (very) redundantly – good for reliability**

# Backups

## Snapshot approach

- At midnight, stop writing into file system
- New writes go into a new file system
  - Mostly pointers to yesterday's data
  - Changes stored in the live file system
    - » Maybe entire files (copy-on-write)
    - » Maybe just new data blocks
- Great for users
  - Old snapshots can be mounted (read-only)
  - Accidentally delete a file? Get it from yesterday!
  - AFS supports a simple version (see “OldFiles”)



# Summary

## **Block-mapping problem**

- Similar to virtual-to-physical mapping for memory
- Large, often-sparse “address” spaces
  - “Holes” not the common case, but not impossible
- Map any “logical address” to any “physical address”
- Key difference: file maps often don't fit in memory

## **“Insert a level of indirection”**

- Multiple file system types on one machine
- Grow your block-allocation map
- ...

# Further Reading

## **Journaling**

- Prabhakaran et al., Analysis and Evolution of Journaling File Systems (USENIX 2005)

## **Something cool which isn't journaling**

- McKusick & Ganger: “Soft Updates: A Technique for Eliminating Most Synchronous Writes in the Fast Filesystem” (USENIX 1999)

**Both papers appear in the “filesystem reliability” book report paper track**