

15-410
“...*RISCy Business*...”

Project 4

Joshua Wise
April 1, 2010

Synchronization

- SIGBOVIK today!
 - 16:00, Rashid Auditorium
 - (That's GHC 4401 to you)
 - (I think)
 - Experience a real live academic conference!

Project 4

- Early announcement this semester
 - *Start designing early!*
- Wait until you have a solid kernel to begin implementing
 - As always, a solid P3 is much more important than a shaky P3 and P4

Project 4: MIPS Port

- Qualitatively different from previous P4s
- Porting an existing Pebbles kernel to MIPS
 - *your* Pebbles kernel!
- Goals: learn about...
 - ...designing for portability
 - ...another architecture
 - ...”wacky” non-x86 machines

Outline

- MIPS, and the RISC way of life
 - Motivations for RISC
 - Implementation specifics...
 - ...and implementation artifacts
 - Extending RISC machines
- Concurrency on MIPS
- VM on MIPS
- Wrap-up

Visit to 1983

- **(or: A Quick Note from 18-447)**
- Performance: a big problem with CPUs
- CISC CPUs, like Intel 80286, slow
 - Many cycles needed per instruction
 - Decoding instructions *difficult*
- What if each instruction could be made simpler?
 - More homogeneous?
 - More *orthogonal*?

Visit to 1983, Part 2

- MIPS project started in 1989
 - John Hennessy, Stanford University
- Experimental technologies for higher performance
- Improve *decode* speed
 - Try to reduce instructions down to one of three forms
- Improve *system* speed
 - Get system frequency up

Peace through Superior Frequency

- If we can simplify things...
 - ...can we get the system down to one clock per instruction?

Peace through Superior Frequency

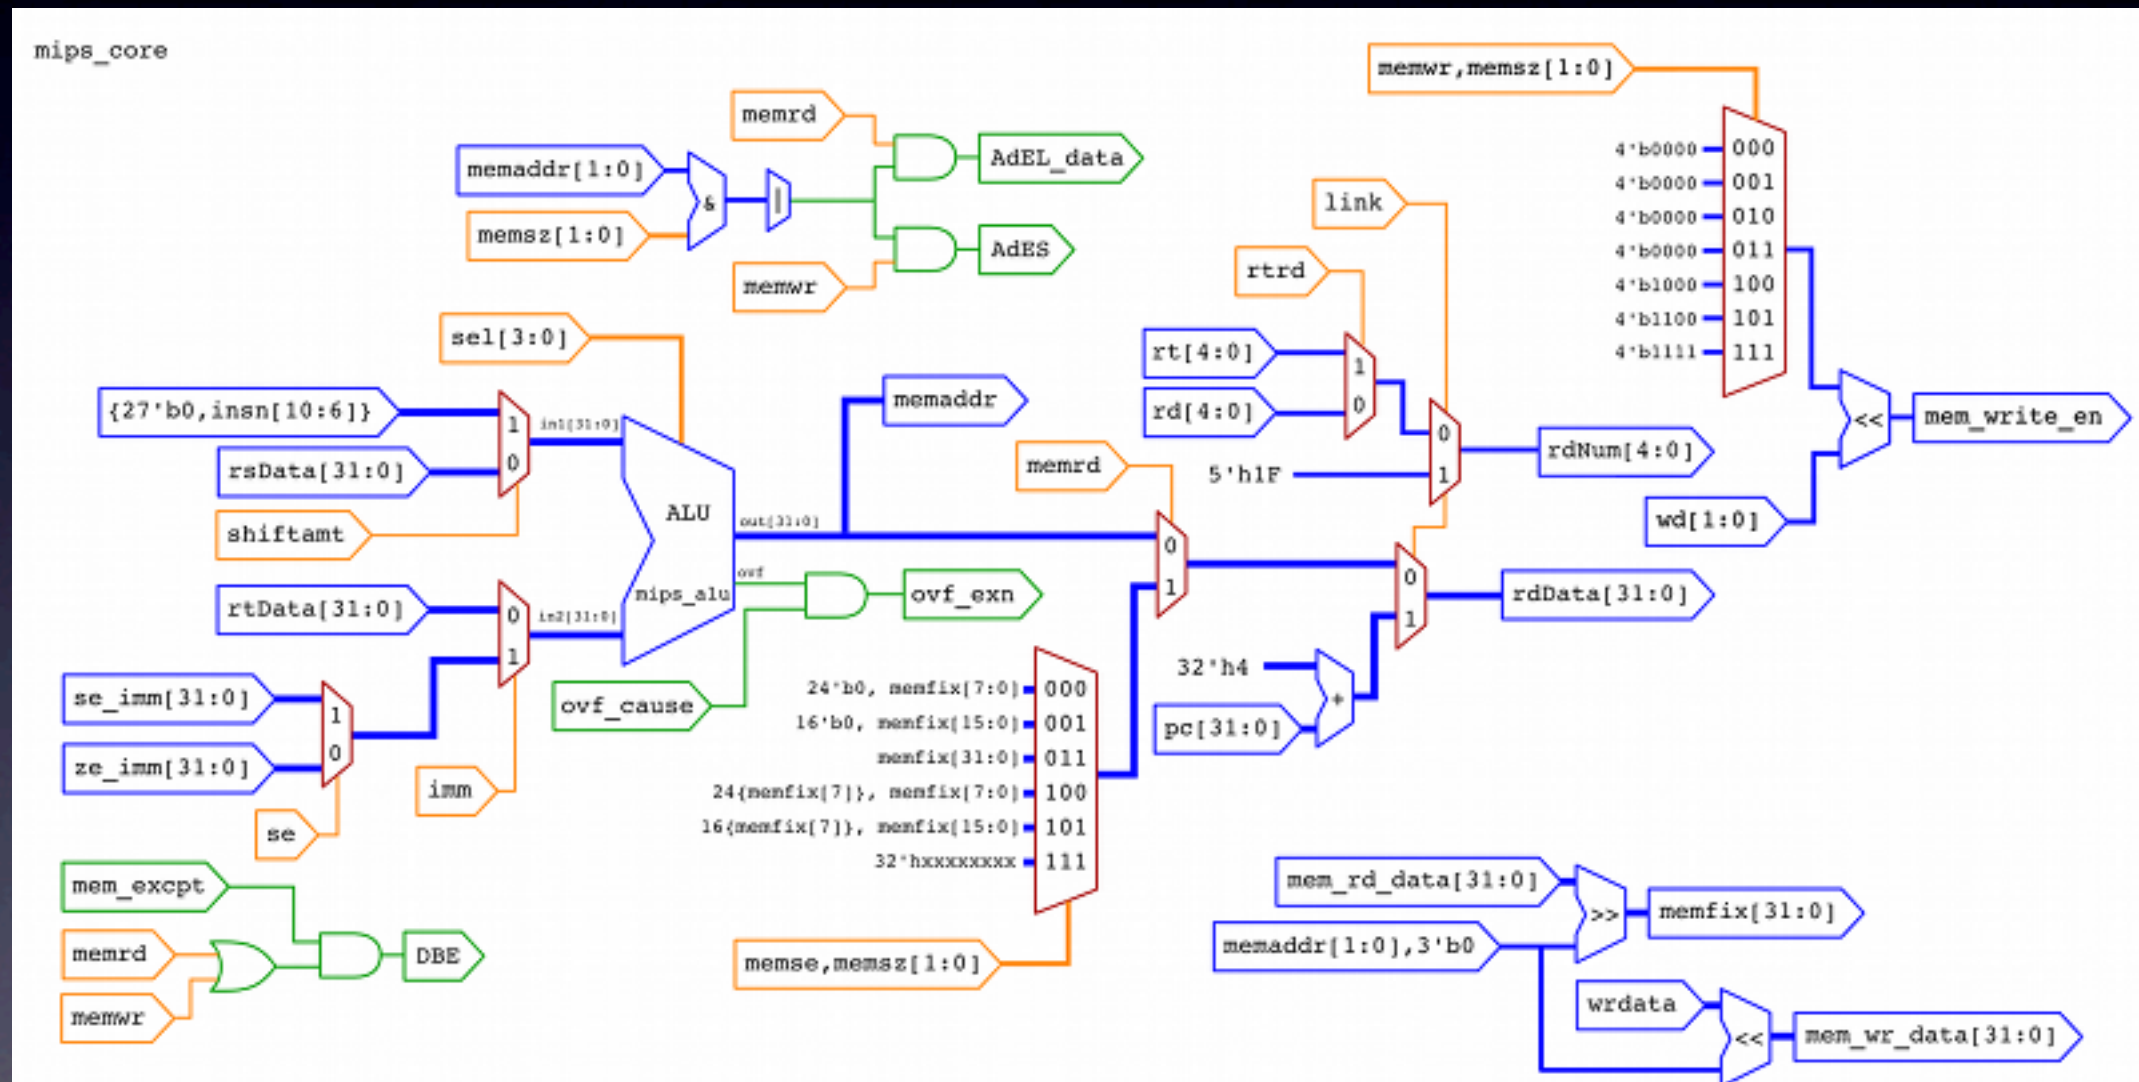


Image Credit: myself, Christopher Lu, Jacob Potter

Actually Superior Frequency

- We got it done in one clock...
 - ...but that is *one slow clock*.
 - *Critical path*: longest path needed for data from one clock to next
- Pipelining
 - What if we *average* one clock per instruction?
 - Each clock, some *part* of an instruction completed
 - Many instructions *in flight*
 - Every clock, one instruction comes out the end

Actually Superior Frequency




Image Credit: Wikipedia user Inductiveload

Actually Superior Frequency

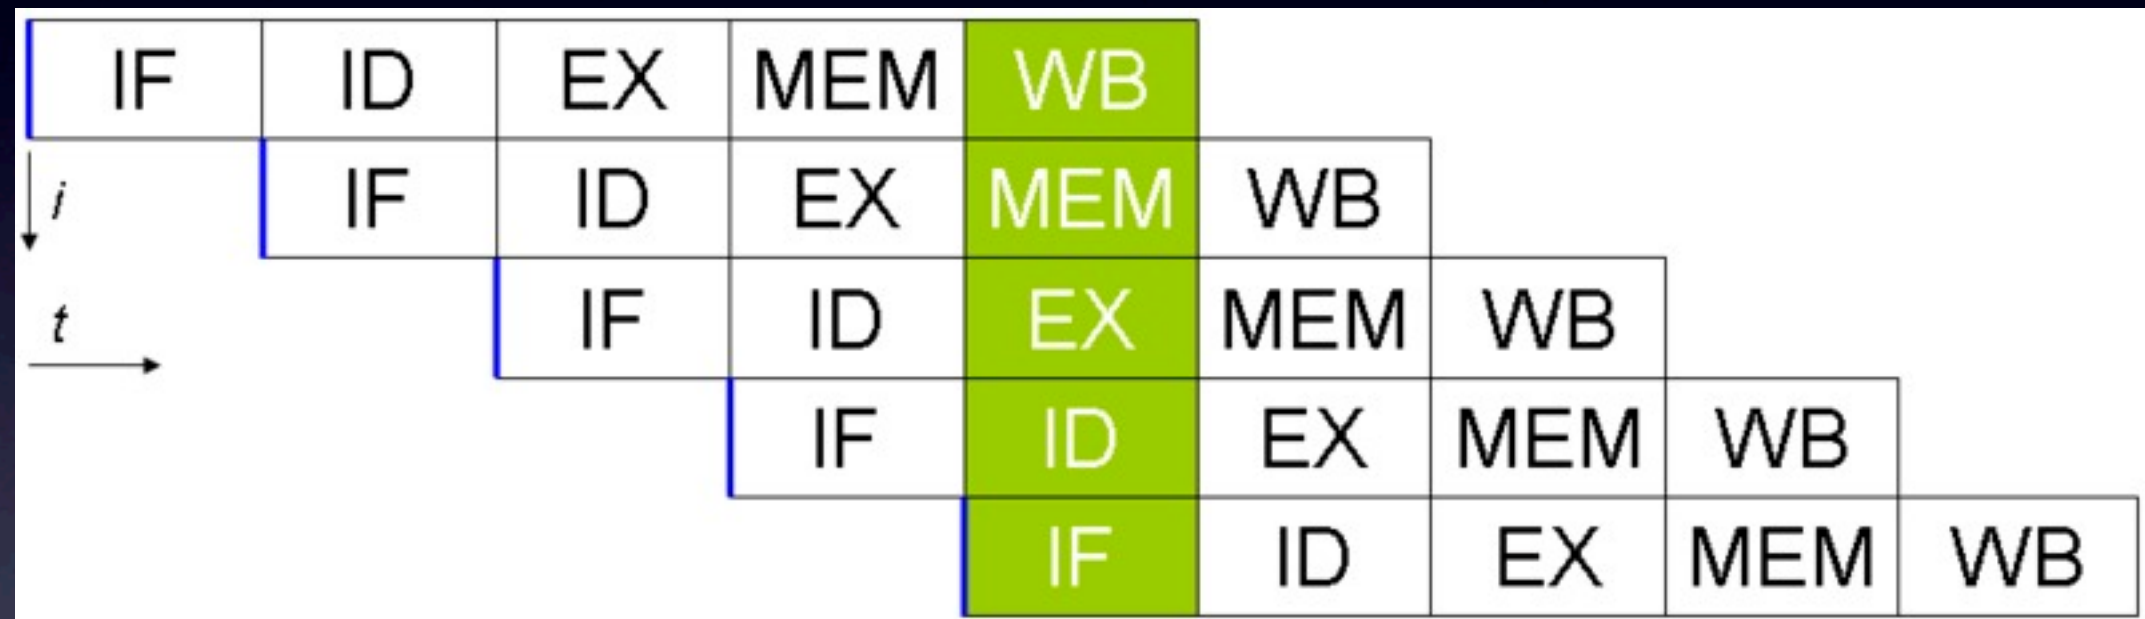


Image Credit: Wikipedia user Poil

MIPS Specifics

- `</ece>`
 - Want to know more? Consider 18-447
 - (My *other* class this semester!)
- All instructions are four bytes
- 32 registers
 - ...almost -- register \$0 is always zero
- Program counter is not a register!
- “Load-store” machine
 - Cannot directly manipulate memory words
- Other than that, *it's just a computer*

A Look at MIPS Instructions

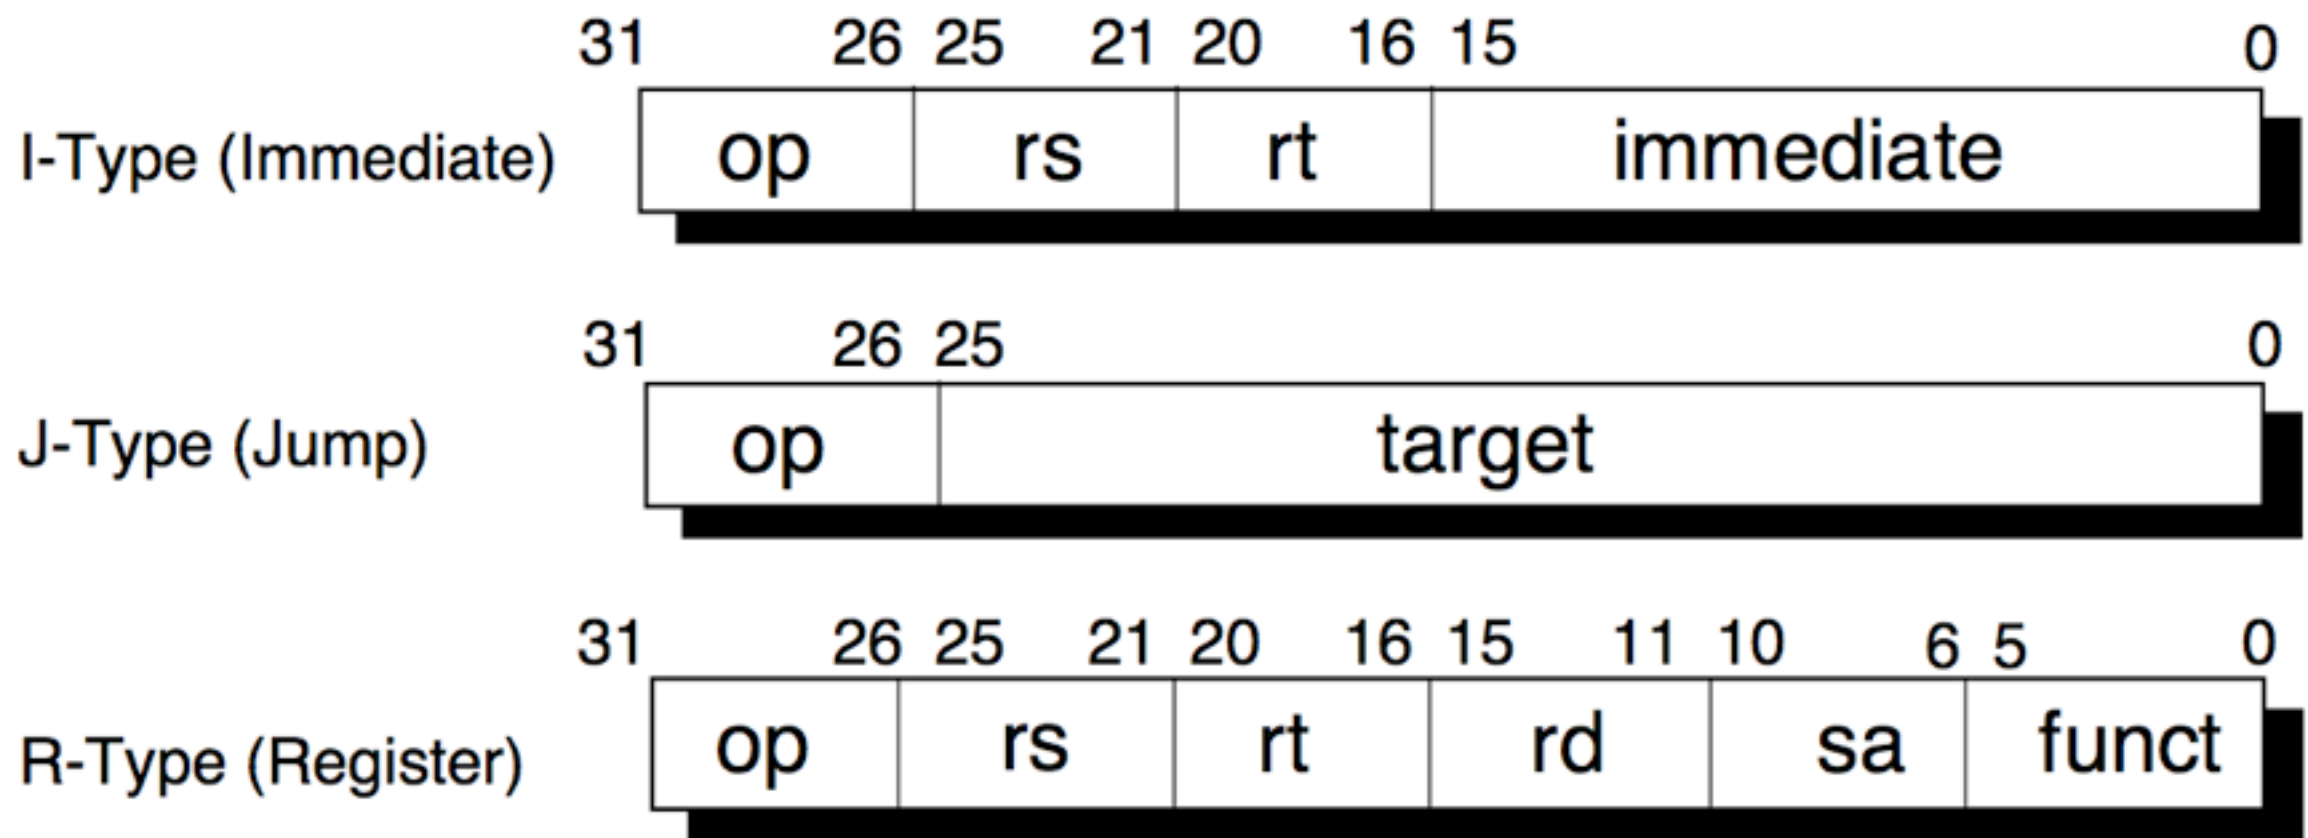


Figure 1-3 CPU Instruction Formats

Image Credit: MIPS R4000 Microprocessor User's Manual, Joe Heinrich

A Look at MIPS Instructions

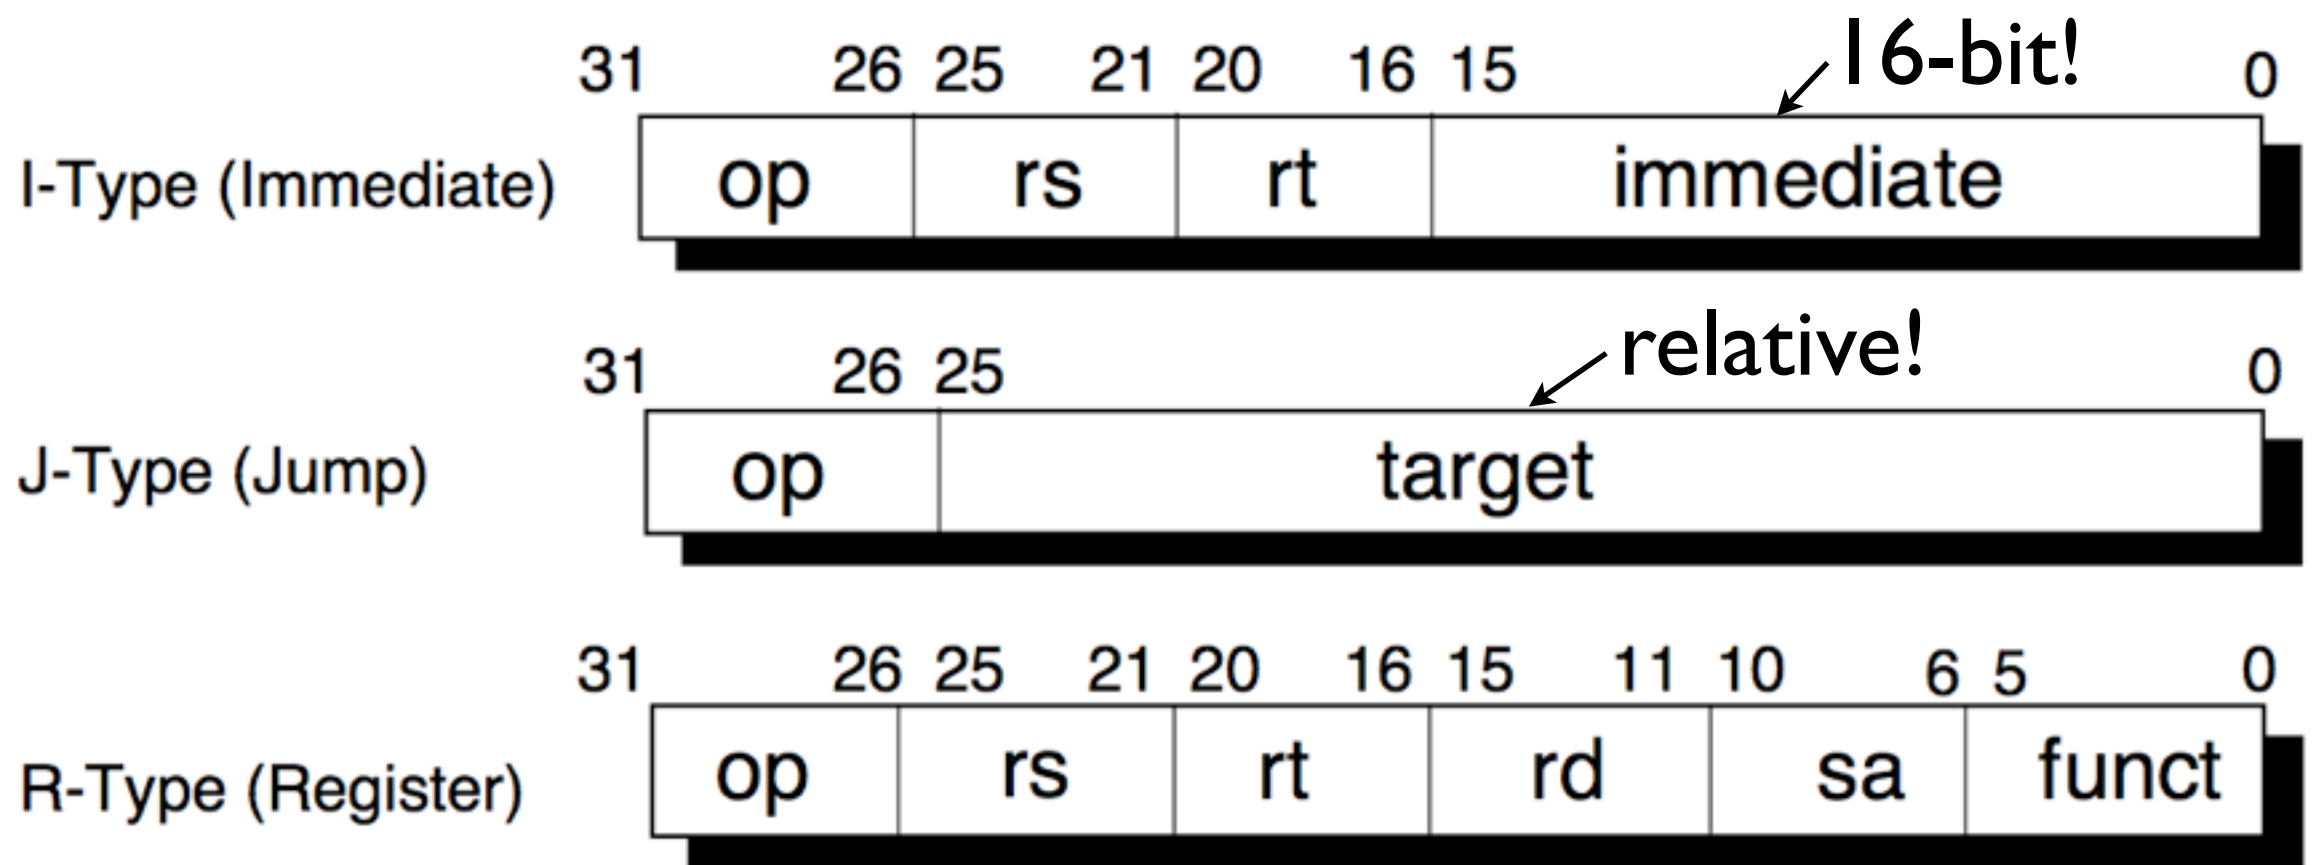
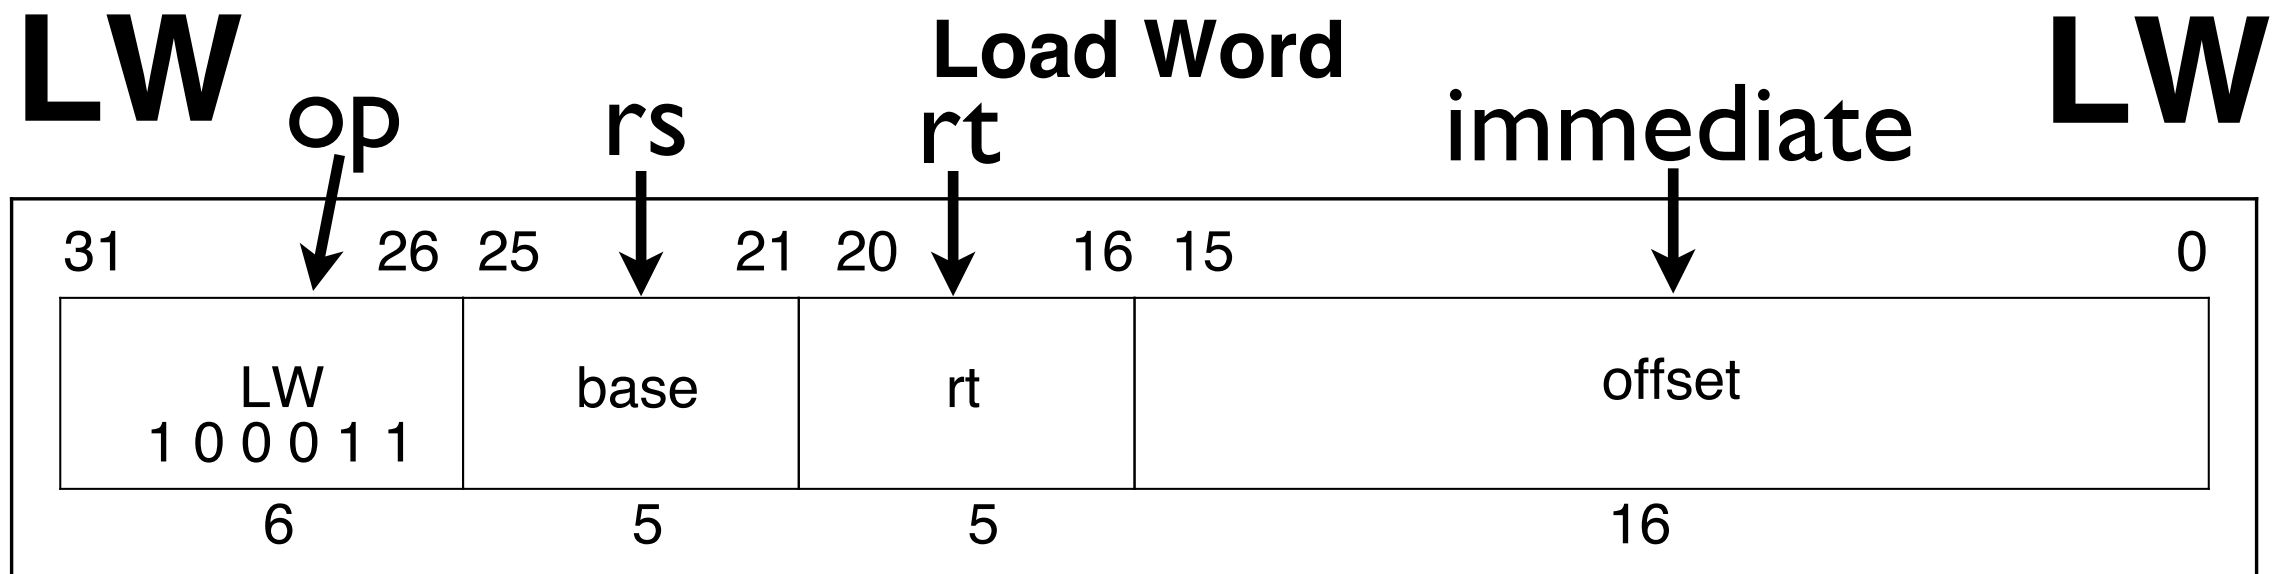


Figure 1-3 CPU Instruction Formats

Image Credit: MIPS R4000 Microprocessor User's Manual, Joe Heinrich

A Look at MIPS Instructions



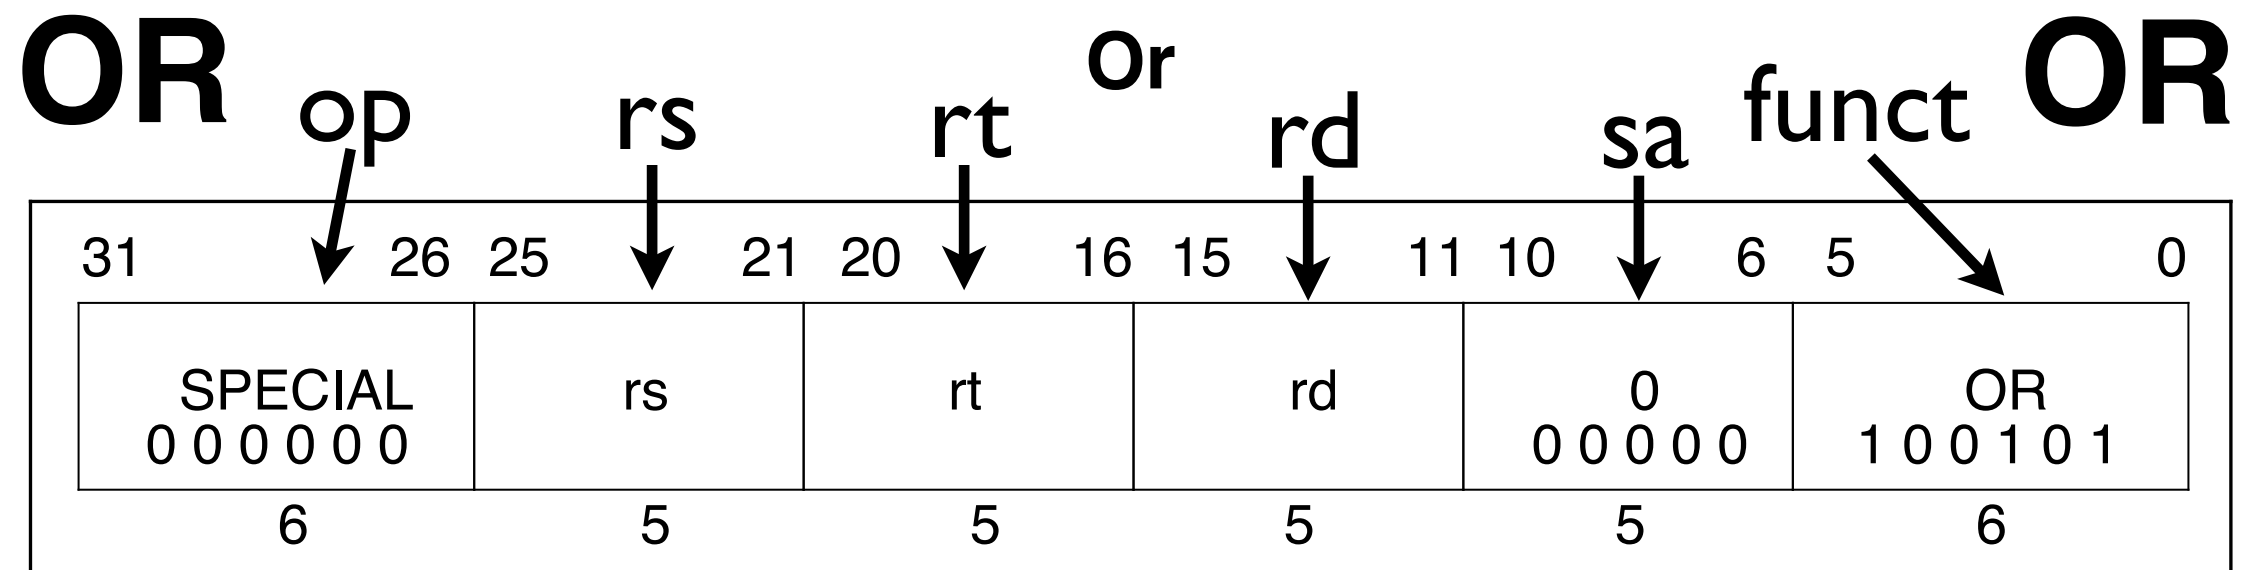
Format:

LW rt, offset(base)

I-type!

Image Credit: MIPS R4000 Microprocessor User's Manual, Joe Heinrich

A Look at MIPS Instructions



Format:

OR rd, rs, rt

R-type!

Image Credit: MIPS R4000 Microprocessor User's Manual, Joe Heinrich

What's interesting about that?

- *They all look the same!*
 - There are only **three** types.
- What's the instruction for a 32-bit immediate?
 - Where do the bits go?
 - ```
lui $v0, 0x1234 // v0 <- 0x12340000
addi $v0, $v0, 0x5678 // v0 <- v0 + 0x5678
 // (v0 = 0x12345678)
```

# A Look at Some Code

```
void swap(int *xp, int *yp)
{
 int t0 = *xp;
 int t1 = *yp;
 *xp = t1;
 *yp = t0;
}
```

# A Look at Some Code

## x86

```
// src, dest
swap:
 pushl %ebp
 movl %esp, %ebp
 pushl %ebx

 movl 12(%ebp), %ecx
 movl 8(%ebp), %edx
 movl (%ecx), %eax
 movl (%edx), %ebx
 movl %eax, (%edx)
 movl %ebx, (%ecx)

 movl -4(%ebp), %ebx
 movl %ebp, %esp
 popl %ebp
 ret
```

## MIPS

```
// dest, src
swap:
 addiu $sp, $sp, -4
 sw $fp, 0($sp)
 move $fp, $sp

 // arguments passed
 // in registers
 lw $t0, 0($a0)
 lw $t1, 0($a1)
 sw $t0, 0($a1)
 sw $t1, 0($a0)

 move $sp, $fp
 lw $fp, 0($sp)
 jr $ra
 addiu $sp, $sp, 4
```



# MIPS Calling Convention

- Since you have so many registers...
  - why not use them all?
- These slides are *not authoritative*
  - A few of the more interesting highlights
  - Real reference in handout some time next week
- Many arguments are in registers, not stacked
  - \$a0 through \$a7

# MIPS Calling Convention

- One register reserved for the assembler -- `$at`
  - MIPS instructions small; assembler provides macroinstructions!
  - i.e., `bgt $rs, $rt, label`
  - becomes:  
`slt $at, $rs, $rt`  
`bne $at, $zero, label`
- Two registers reserved for the *kernel*
  - Have to compute an address to save registers to!
  - `$k0` and `$k1` -- may change *while a user mode program is running*

# MIPS Calling Convention

- One “*global pointer*”
  - Offset accesses for global data
  - “Partial pointer”
  - Impossible: `lw $rt, myPointer($0)`
  - Replacement: `lw $rt, gpBase-myPointer($gp)`
  - Preserved throughout lifetime of program
- Other more mundane bits (caller-save, callee-save)
  - See handout, or reference materials on web





# Coprocessors

- If instructions are simple...
  - ...then how do we do hard things?
- There *is no* “**CPUID**” instruction
- **CPUID** lives on “coprocessor zero”
- *Don’t be fooled by the name!*
  - Coprocessors need not do any actual processing
- Other coprocessors include:
  - FPU
  - Multiply/divide (on some MIPSen)
- Access with special instructions

# Magic Coprocessor Zero

- One coprocessor is more magical than the rest
- Present on *every single MIPS*
- Contains controls for:
  - Cache control
  - System identification (“CPUID”)
  - Exception/interrupt control
  - Timers, memory control, ...
  - *and virtual memory!*
- VM, PIC, PIT, and northbridge rolled into one

# set\_cr3?

## x86

```
// src, dest
set_cr3:
 movl 4(%esp), %eax

 movl %eax, %cr3

 ret
```

## MIPS

```
// dest, src
set_context:
 // argument already
 // in $v0

 jr $ra
 mtc0 $v0, C0_CONTEXT
```

**C0\_CONTEXT** is an *optimization*:  
“Where you should look” on a fault



# Concurrency on Earth

- We know how to do atomic operations on x86
- `lock xchg`, discussed in Synch #2
- Implements *one and only one* concurrency mechanism
- Write and read characteristics defined
- ```
atomic int32 xchg(int32 *mem, int32 new) {  
    register int old;  
    old = *mem;  
    *mem = new;  
    return old;  
}
```

Earth Technology is Behind

- What about something else?
 - xchg... conditionally?
 - Atomic addition?
 - *Add and get old value?*
 - *Atomically two's complement memory???*
- CISC solution

Earth Technology is Behind

- What about something else?
 - xchg... conditionally?
 - Atomic addition?
 - *Add and get old value?*
 - *Atomically two's complement memory???*
- CISC solution: new instruction for each
- **CMPXCHG, CMPXCHG8B, LOCK ADD, LOCK XADD, LOCK NOT**
 - (LOCK? What's that?)

The Problem with x86 is...

- Adding more instructions is not very RISC-like
- Adds complexity to decode stage
- *Reduces flexibility of system*
- Atomic operations do not fit in pipeline!
 - Two memory operations
 - Waiting is **not** an option

A RISC Solution

- Can we factor out a *common element* from concurrent operations?
 - Do something...

A RISC Solution

- Can we factor out a *common element* from concurrent operations?
 - Do something...
 - ...*only if nobody got in our way.*

A RISC Solution

- Load linked (LL)
 - Load memory, and set the link flag
- Link flag
 - *Cleared* when somebody else writes to the same address (cache shutdown, or context switch)
- Store conditional (SC)
 - Store memory *only if the link flag is set*
 - Informs programmer if SC failed

Rewriting mutex_lock

- LL/SC can be used to implement *any* of those x86 instructions
- Spin-wait in MIPS assembly
 - ```
.1:
 ll $t0, lock_avail($gp)
 bz $t0, .1 // branch if zero
 addi $t0, $zero, 0 // i.e., $t0 = 0
 sc $t0, lock_avail($gp) // stores 1 if successful
 bz $t0, .1 // have to retry if 0
 jr $ra // like 'ret'
```
- No special operation needed for unlock!

# VM on Earth

- On x86, action centered around page table
- One instruction could cause:
  - ...



# VM on Earth

- On x86, action centered around page table
- One instruction could cause:
  - TLB lookup, and miss
  - PDBR (%cr3) read
  - Page directory address calculation
  - Page directory dword read
  - Page table address calculation
  - Page table dword read
  - TLB eviction decision, evict, and store
  - TLB lookup *again*
  - Memory read for actual instruction
- That is a **LOT** of logic!
  - ... and just to handle *one* instruction!

# Earthly Ties

- This is in keeping with the CISC way
  - Big complicated task? *Let the hardware deal with it!*
- There is only *one way* to organize pages
  - Want extra bits? *Too bad!*
- Logic required to deal with this is immense
- “Locked in” to one page table format
  - Reserved bits, and PAE/PGE
- “Encapsulate your indefensible code”!

# VM, MIPS-Stylin'


- Once again, what is the *common element* to determine if an access is valid?
- Our friend we love to hate...



# VM, MIPS-Stylin'

- Once again, what is the *common element* to determine if an access is valid?
  - Our friend we love to hate...
  - *the TLB!*
- MIPS abstracts virtual memory *only to the TLB level*
- Kernel memory “locked” in the upper half of the address space
  - **kseg**
- Let's trace the same case again

# Address space



# Martian TLB Misses

- TLB lookup for address misses
- Fix needed, *CPU takes exception!*
  - CPU vectors to exception handler
  - EH reads faulting address from cop0
  - *Exception handler* picks “random” entry to evict
    - **TLBWR**
  - *Exception handler* reads internal data structures to “fix” fault
  - *Exception handler* traps to kernel if needed
  - *Exception handler* refills TLB
  - **ERET** back to retry instruction
- TLB lookup hits; life proceeds



# Sample Exception Handler

- The hardware does help
  - If you set up C0\_CONTEXT, it will fill in where to read from
- Best case:
  - ```
mfc0 k1, C0_CONTEXT
lw    k0, 0(k1)
lw    k1, 8(k1)
mtc0 k0, C0_ENTRYLO0
mtc0 k1, C0_ENTRYLO1
ssnop
ssnop
ssnop
tlbwr
eret
```

```
// Otherwise, the nop
// would get parallelized.
// I'm not joking.
```

TLB Exception Notes


- TLB exception must be *fast*!
 - On MIPS R4K (ours), TLB exceptions usually ~25 instructions
 - Still competitive with x86 TLB miss!
- TLB exception generally follows the same procedure as x86 hardware would
- Exn handler usually not portable
- Looking for hints writing it?
 - `linux/arch/mips/mm/tlbex.c` -- runtime generated
 - `build_r4000_tlb_refill_handler`
 - Yours probably won't be runtime-generated

Exceptions on MIPS

- Exceptions on MIPS also a whole lot less complicated
- MIPS
 - No IDTR, no IDT
 - Exception comes along (TLBL exception!)
 - Processor *immediately* vectors to *fixed* location
- Exceptions *faster*
 - Important if you have a software TLB!
- Exception reason indicated by cop0 **CAUSE** register
- We will provide infrastructure to help set these up

Mechanics

- Simics doesn't support our MIPS platform very well
- *Two methods for you to run your code*
 - MIPS Malta
 - Generously donated to Carnegie Mellon when SiCortex went under
 - Debugging using EJTAG probe
 - Instructions soon



Mechanics

- *Two methods for you to run your code*
 - `tmips`: developed by Thomas Tuttle for 18-447
 - Make sure to thank him if you see him!
 - Software simulator
 - Bored before P3 ends?
 - <http://www.github.com/ttuttle/tmips>

Mechanics

- *Three methods for you to run your code?*
 - *This very MIPS processor is being implemented by 18-447 students right now!*
 - Discussed today: MIPS is very simple
 - Simple enough for an *intro comp arch class* to implement almost in its entirety
 - 18-447 MIPS is missing a *few* CP0 features...
 - ...but could be trivially extended to support Pebbles
- Excited? Go take 18-447 :-)

Wrap-up

- We talked about:
 - Motivations for RISC
 - Implementation specifics for MIPS
 - Concurrency on MIPS
 - VM on MIPS
 - Exceptions
- P4 handout to be released on day mentioned on 15-410 web site
- Don't spend too much time worrying about MIPS until you have a *solid* kernel!
- Some of these ideas may be on the final!