

CMSC 412

Operating Systems

Fall 2004

Professor Michael Hicks

Welcome!

- Read the Syllabus
 - read the warning about the size of the project
 - make sure you get the 6th edition (or later) of the book
- Discussion Sections (start Wednesday)
 - focus on the project
 - meet only once a week
 - Probably we will have three times total per week; you can attend the section of your choice. See the web page.

Projects

- The best way to understand is by doing
 - So, we will build an operating system that works on bare hardware by extending the GeekOS academic operating system. 6 projects total.
- Project #0 Handout (posted August 31)
 - It will be due late next week.
 - It will get you familiar with the simulator and the Cyclone programming language, which you can use to write some or all of the project.

What is Cyclone? Glad you asked ...

Cyclone

- Based on C
 - Low-level control over memory management, data representation, and access to the machine
- But type-safe!
 - Rules out many hard-to-find bugs and security holes
 - Buffer overflows
 - Dangling pointers
 - These bugs have killed many projects

Cyclone in this class

- You will be required to use it for project 0, to get familiar with it
- From then on it is optional
 - Could very well make your project more reliable, and prevent many bugs
 - But, it's a research compiler, and it will have bugs itself, and may have cryptic error messages
 - We will do our best to overcome these issues with you. If a major problem in the compiler inhibits a good-faith attempt to complete the project, you will not be penalized.

Cyclone in this class

- Completing at least one project using Cyclone (from project 1 to 6) will net some extra credit
 - You must include a writeup of your experience, including what features you used, what difficulties you ran into, and suggestions for improvement.
- See <http://www.cs.umd.edu/projects/cyclone/>

Class Grades Server

<http://grades.cs.umd.edu>

- Get your LinuxLab account from here
 - CS computing cluster. Projects must work and be submitted on these machines.
- Complete grade information
- Interface for requesting regrades on exams and projects (with deadline!)

Course Material

- Reading
 - Chapter 1
 - Chapter 2 (for Wednesday)
- Coordination with other section
 - lectures and exams may differ
 - projects will be the same

Why Study OSs?

- Understand computer systems
 - From the hardware to applications
 - Helpful for understanding performance, security, reliability, and other issues

Why Study OSs?

- Understand principles of abstraction
 - OSs are large and complex. How do we manage this complexity?
 - Abstraction:
 - Break each piece into self-contained chunks
 - Plug each piece together
 - Different views of service from view of the provider and view of the user
 - These principles will serve you in other software systems projects
 - E.g., thinking about concurrency

Why Study OSs?

- Understand tradeoffs of system design
 - Many types of users (too many!)
 - real-time, desktop, server, etc...
 - Many possible models and abstractions
 - OS researchers are 'abstraction merchants'
 - There is no perfect OS!

Why Study OSs?

- It's fun!
 - the details are interesting (I think so!)

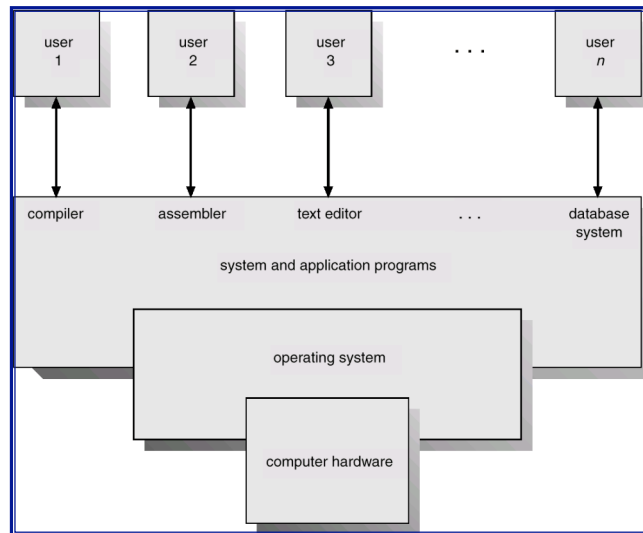
What is an Operating System?

- Resource Manager
 - Resources include: CPU, memory, disk, network
 - OS allocates and de-allocates these resources
- Virtual Machine
 - provides an abstraction of a larger (or just different) machine
 - Examples:
 - Virtual memory - looks like more memory
 - Java - pseudo machine that looks like a stack machine
 - IBM VM - a complete virtual machine (can boot multiple copies of an OS on it)
- Multiplexor
 - allows sharing of resources and protection
 - motivation is cost: consider a \$40M supercomputer

What is an OS (cont)?

- Provider of Services
 - includes most of the things in the above definition
 - provide “common” subroutines for the programmer
 - windowing systems
 - memory management
- The software that is always loaded/running
 - generally refers to the OS *kernel*.
 - small protected piece of software
- All of these definitions are correct
 - **but** not all operating have all of these features

Abstract View of System



Operating System Cousins

- **Hardware**
 - OS is managing hardware resources so needs to know about the ugly details of the hardware
 - interrupt vectors
 - page tables
 - I/O registers
 - some features can be implemented either in hardware or the OS
- **Languages**
 - can you write an OS in any language?
 - No: need to be able to explicitly layout data structures to match hardware

OS Cousins (cont)

- Language Runtime systems
 - memory management requirements
 - explicit heap management
 - garbage collection
 - stack layout
 - concurrency and synchronization
 - calling convention (how are parameters passed)
- Data Structure and Algorithms
 - efficient access to information in an OS
 - for most things need linear time and space
 - for many things want log or constant time

OS Examples: Batch system

- Mainframes
 - Automate running of user jobs
 - Users submit jobs, OS schedules them
 - Called a **batch** system
 - OS referred to as a **resident monitor**

Multiprogramming

- Many jobs active “at once”
 - Each job no longer runs to completion before the next can start
 - Many jobs resident in memory at once; the CPU is shared between them
- Job scheduling
 - which jobs to load into memory

Multiprogramming OSs

- Must manage resources
 - Memory
 - I/O devices
 - CPU
 - Want to be fair and efficient
- CPU scheduling
 - which memory-resident program to run
- Protection
 - Prevent errant job from tainting results of another

Multitasking (Timesharing)

- Like multiprogramming, but **interactive**
 - Switching between tasks is very fast
- **Preemption**: interrupting a job to transfer control to another job
 - **Quantum**: the time slice allocated to a process before it's preempted. Modern OS typically 50 ms.
- **Interactivity** requires more support
 - Filesystem, virtual memory, synchronization, deadlock avoidance

Desktop System

- Catered to a single user
- **Goal**: flexibility and responsiveness
 - Support interactive I/O devices
 - May not need to be as efficient as multi-user OS to be more responsive
 - May be less concerned with protection

Parallel System

- Single computer system of several tightly-coupled CPUs
 - All share memory, clock, devices
- Goals:
 - Better value
 - Increased performance
 - Lower cost
 - Increased reliability
- Two flavors
 - symmetric and asymmetric

Distributed System

- Many machines joined by a network
 - Loosely coupled: each processor has its own memory; communication via the network
 - OS often implements network protocols
- Goals
 - Value
 - Share distributed resources
- Variants: client-server, peer-to-peer

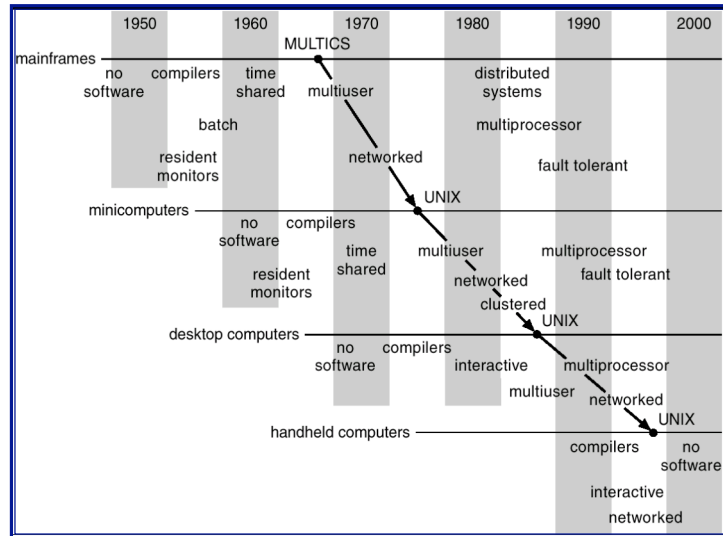
Real-time Systems

- Control device in a dedicated application
 - Medical imaging, industrial control, etc.
- Well-defined fixed time constraints
- Variants
 - **Hard**: can never miss a deadline
 - **Soft**: less stringent. Thus permits unpredictable elements like caches, disks, etc.

Handheld Systems

- PDAs, cell phones, etc.
- Issues
 - Limited memory
 - Slow (power-constrained) processors
 - Small display

Migration of OS features



OS Goals: Usability

- **Robustness**
 - accept all valid input
 - detect and gracefully handle all invalid input
 - should not be possible to crash the OS
- **Consistency**
 - same operation should mean the same thing
 - read from a file or a network should look the same
 - a “-” flag should be the same in different commands
 - conventions
 - define the convention
 - follow the convention when adding new items

Usability Goals (cont)

- Proportionality
 - simple, common cases are easy and fast
 - good default values
 - complex, rare cases are possible but more complex and slower
 - “rm *” should give a warning
 - formatting the disk should not be on the desktop next to the trash can

Cost Goals

- Good Algorithms
 - time/space tradeoff are important
 - use special hardware where needed
 - smart disk controllers, memory protection
- Low maintenance cost
 - should not require constant attention
- Maintainability
 - most of cost of software is in maintenance so make it easy to maintain the software base

Adaptability Goals

- Tailored to the environment
 - server vs. workstation
 - multi-media vs. data entry
- Changes over time
 - added memory
 - new devices
- Extensible
 - third parties can add new features
 - database vendors often need custom features
 - end customers can extend the system
 - new devices
 - new policies