

CMSC 412


Fall 2004

Computer and Operating System
Structures

Announcements

- Project #0
 - Due Friday
 - Submission program is now working
- Project #1
 - Much harder than Project #0!
 - Posted tomorrow
- Reading
 - Chapter 2,3
 - Chapter 3,4 (for Monday)

Computer Systems



I/O Subsystem

- Many different types of devices
 - disks
 - networks
 - displays
 - mouse
 - keyboard
- Each has different peak performance
 - bandwidth
 - rate at which data can be moved
 - latency
 - time from request to first data back

Impacts System Structure

- Different Buses (not just one!)
 - Processor Bus (on chip)
 - Many Gigabytes/sec
 - Memory Bus (on processor board)
 - ~1-2 Gigabytes/sec
 - I/O Bus (PCI, MCA)
 - ~100 Megabytes/sec
 - Device Bus (SCSI, USB)
 - tens of Megabytes/sec

Impacts System Structure

- Different ways to transmit data
 - Interrupt-driven I/O
 - CPU directs how data is transmitted to and from the device
 - User interrupts CPU to tell data to be transmitted
 - Device interrupts CPU when task is completed or data is available
 - Direct Memory Access (DMA)
 - CPU directs transmission at a much more coarse-grained level

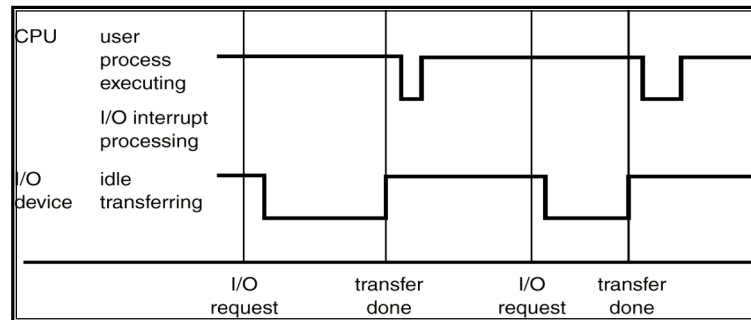
I/O Interrupts

- Indicate an event has occurred
 - can be caused by hardware devices
 - indicates data present or hardware available
 - can be caused by software
 - *system call* (or *trap*)
- Interrupt the CPU to execute a *handler*
 - saves state about what was happening
 - returns where it left off when finished

Servicing Interrupts

- Need to know what device interrupted
 - could ask each device (slow!)
 - instead use an *interrupt vector*
 - array of pointers to functions to handle a specific interrupt
- What happens if an interrupt arrives while we are in an interrupt handler?
 - OS may wish to *disable interrupts* until handling is complete.

Servicing Interrupts



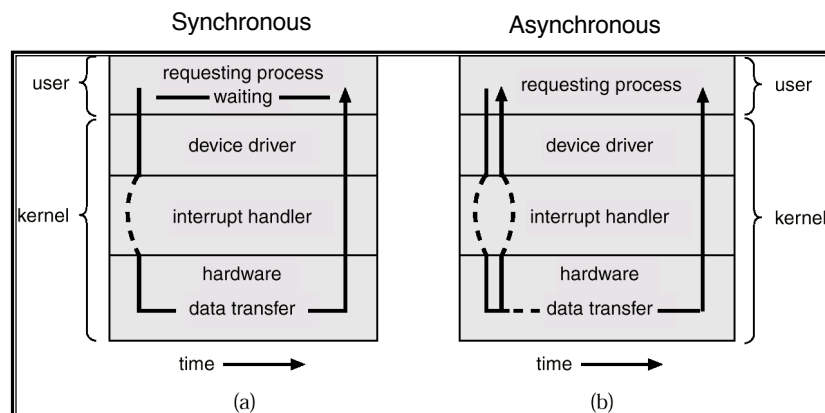
Synchronous I/O

- After I/O starts, control returns to user program only upon I/O completion.
 - *Wait* instruction or loop idles the CPU until the next interrupt
 - At most one I/O request is outstanding at a time; no simultaneous I/O processing.
- Modern OSs do not perform synchronous I/O
 - Not compatible with multitasking
 - May provide the illusion of synchronous I/O to user programs, however

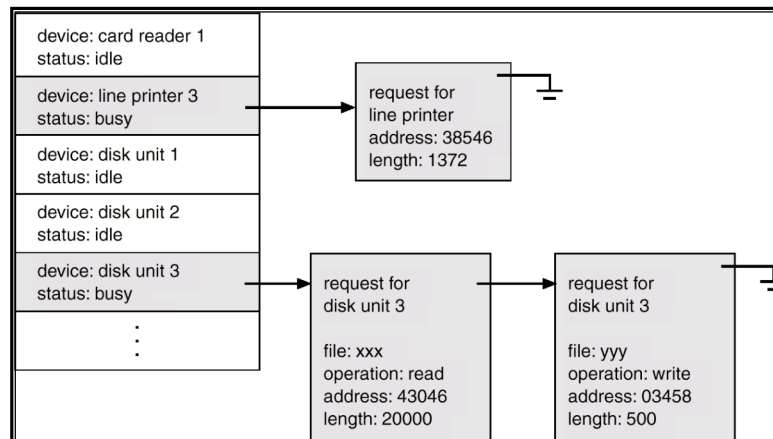
Asynchronous I/O

- After I/O starts, control returns to (some) user program without waiting for I/O completion.
 - *Device-status table* contains entry for each I/O device indicating its type, address, and state.
 - Operating system indexes into I/O device table to determine device status and to modify table entry to process interrupt.

Two I/O Methods



Device-status Table



Direct Memory Access (DMA)

- Used for high-speed I/O devices able to transmit information at close to memory speeds.
- One interrupt is generated per *block*, rather than the one interrupt per *byte*.
- Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention.

System Calls

- A software-generated interrupt
 - a.k.a. *trap*
- Provide the interface between application programs and the OS kernel
- Are like procedure calls
 - take parameters
 - calling routine waits for response

System Call Mechanism

- Use numbers to indicate what call is made
- Parameters stored in registers or the stack
- Why do we use system call numbers rather than directly calling a kernel subroutine?
 - permits changing the size and location of system call implementations without having to re-link application programs

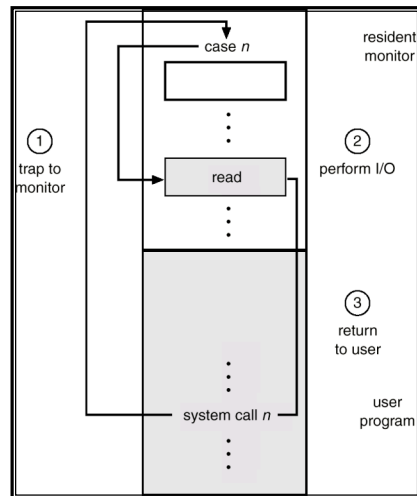
GeekOS and x86

- Intel system call instruction
 - `int n` where *n* is the interrupt vector #
 - “call kernel routine *n*”
 - vectors 0-31 reserved
 - Page fault, segmentation violation, etc.
- GeekOS
 - All system calls set *n* as 90
 - System call number stored in *eax* register

Types of System Calls

- File Related
 - open, create, read, write, close, delete
 - get or set file attributes
- Informational
 - get time
 - set system data (OS parameters)
 - get process information (id, time used)
- Communication Related
 - establish a connection; terminate a connection
 - send, receive messages
- Process control
 - create/terminate a process (including self)

Use of a System Call for I/O



Why use system calls at all?

- Why not “link” application programs against the kernel and call kernel routines directly?

Protection


- Need to protect OS from user programs and user programs from each other
 - Don't want a bug in a user program to crash the whole machine (as in earlier OSs, like MS-DOS, MacOS, Windows 3.1, and others)
- Hardware resources of interest
 - Memory, I/O devices, CPU

Dual-Mode Operation

- Provide hardware support to differentiate between at least two modes of operations.
 1. *User mode* - execution for a user.
 2. *Monitor mode* (also *kernel mode* or *system mode*) - execution done on behalf of operating system.
- Operations in user mode a subset of those allowed in monitor mode
 - **Privileged instructions** only in monitor

Dual-Mode Operation

- *Mode bit* added to computer hardware to indicate the current mode: monitor (0) or user (1). X86 actually has 4 modes (2 bits).
- When an interrupt or fault occurs hardware switches to monitor mode.



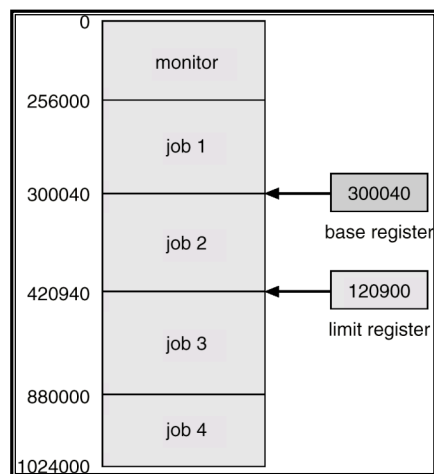
I/O Protection

- System call mechanism prevents user mode programs from accessing devices directly.
 - All I/O instructions are privileged
- But what if user program can overwrite interrupt handler with its own code?

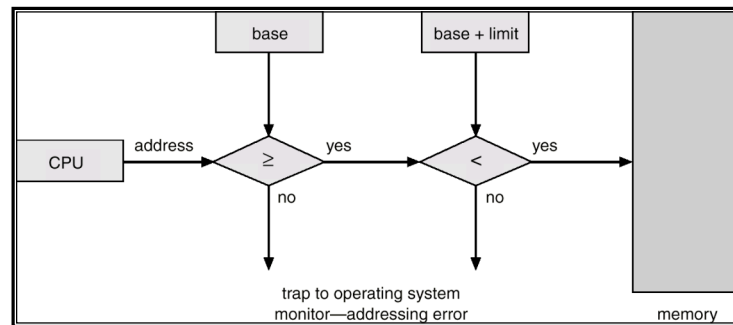
Memory Protection

- Must provide memory protection at least for the interrupt vector and the interrupt service routines.
- In order to have memory protection, add two registers that determine the range of legal addresses a program may access:
 - **Base register** - holds the smallest legal physical memory address.
 - **Limit register** - contains the size of the range
- Memory outside the defined range is protected.

Base and Limit Registers



Hardware Address Protection



Changing the base and limit registers are privileged operations

CPU Protection

- *Timer* - interrupts computer after specified period to ensure operating system maintains control.
 - Timer is decremented every clock tick.
 - When timer reaches the value 0, an interrupt occurs.
- Commonly used to implement time sharing.
- Load-timer is a privileged instruction.

Storage Structure

- Main memory - only large storage media that the CPU can access directly.
- Secondary storage - extension of main memory that provides large *nonvolatile* storage capacity.


Disks

- Several types:
 - **Hard Disks** - rigid surface with magnetic coating
 - **Floppy disks** - flexible surface with magnetic coating
 - **Optical (CDs and DVDs)** - read only, write once, multi-write
- The *disk controller* determines the logical interaction between the device and the computer.

Hard Disks

- Collection of **platters**
- Platters contain concentric rings called **tracks**
- Tracks are divided into fixed sized units called **sectors**
- A **cylinder** is a collection of all tracks equidistant from the center of disk
- Current Performance:
 - capacity: megabytes to hundreds of gigabytes
 - throughput: sustained < 10 megabytes/sec
 - latency: milliseconds

Moving-Head Disk Mechanism




Relevant Disk Ops

- Seek
 - Move arm to appropriate cylinder
- Read/Write
 - One sector at a time
 - Must wait until disk rotates to appropriate sector address
- OS must keep track of raw disk addresses to implement filesystems

Storage Hierarchy

- Storage systems organized in hierarchy.
 - Speed
 - Cost
 - Volatility
- *Caching* - copying information into faster storage system; main memory can be viewed as a last *cache* for secondary storage.


Storage-Device Hierarchy




Caching

- Use of high-speed memory to hold recently-accessed data.
- Requires a *cache management* policy.
- Caching introduces another level in storage hierarchy. This requires data that is simultaneously stored in more than one level to be *consistent*.

Migration of A From Disk to Register



Local Area Network Structure



Wide Area Network Structure

