

# CMSC 412

## Fall 2004

### Processes and Threads


## Announcements

- Project #1
  - Due Friday
- Reading
  - Chapter 4 (parts), 5 (parts)
  - Chapter 7 (for Monday)

# Processes


- What is a process?
  - A program in execution
    - Either sequentially or with multiple “threads of control.”
- What’s not a process?
  - A program on a disk - a process is an active object, but a program is just a file

## Computation Abstractions



*A dual-processor computer*

## Processes vs. Threads



*Processes do not  
share data*

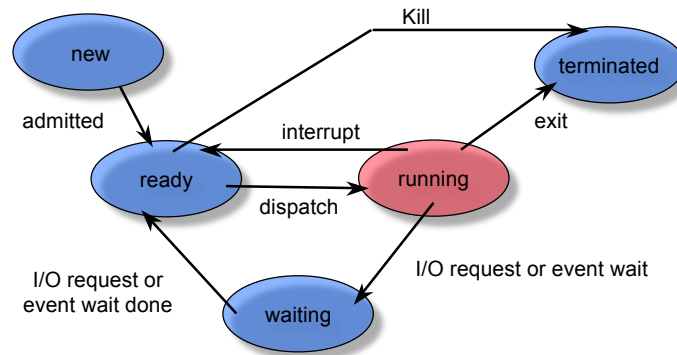
*Threads share data  
within a process*

[More on threads later ...](#)

## Process State

- Processes switch between different states based on internal and external events
- Each process is in exactly one state at a time
- Typical States of Processes (varies with OS)
  - New: just been created
  - Running: instructions are being executed
    - only one process per processor may be running
  - Waiting: waiting for an event to occur
    - examples: I/O events, signals
  - Ready: waiting to be assigned the CPU
  - Terminated: finished execution

## Process State Transitions



## Components of a Process

- Memory Segments
  - Program - often called the text segment
  - Data - global variables
  - Stack - contains activation records
  - Heap - contains dynamically-allocated data
- Processor Registers
  - Control registers
    - program counter - next instruction to execute
    - stack pointer
    - processor status word (from *cmp* instructions)
  - General purpose registers
  - Floating-point registers

## Scheduling

- OS must decide when a process is allowed to run; I.e. it must **schedule** processes
- **Long-term scheduler** (or job scheduler) - selects which processes should be brought into the ready queue.
- **Short-term scheduler** (or CPU scheduler) - selects which process should be executed next and allocates CPU. A.k.a the **dispatcher**.

## Schedulers

- Short-term scheduler is invoked very frequently (milliseconds)  $\Rightarrow$  (must be fast).
- Long-term scheduler is invoked very infrequently (seconds, minutes)  $\Rightarrow$  (may be slow).
- The long-term scheduler controls the *degree of multiprogramming*.
- Processes can be described as either:
  - *I/O-bound process* - spends more time doing I/O than computations, many short CPU bursts.
  - *CPU-bound process* - spends more time doing computations; few very long CPU bursts.

## Scheduling Policy

- How should the scheduler choose a process to run?
- Can simply pick the first item in the queue
  - called **round-robin** scheduling
  - is round-robin scheduling fair?
- Various scheduling criteria
  - Process class (I/O bound vs. CPU bound)
  - Priority
  - Resources consumed
  - Etc.

## Scheduling Implementation

- Use alarm interrupts to switch between processes
  - when time is up, a process is put back on the end of the ready queue
  - frequency of these interrupts (a.k.a. the *quantum*) is an important parameter
    - typically 3-10ms on modern systems
    - need to balance overhead of switching vs. responsiveness

## Context Switch

- When the OS switches a CPU to another process, it is called a **context switch**.
- The system must save the state of the old process and load the saved state for the new process.
- Context-switch time is overhead; the system does no useful work while switching.
  - Total time depends on hardware support.
  - Writing context-switch routines almost always requires some assembly language

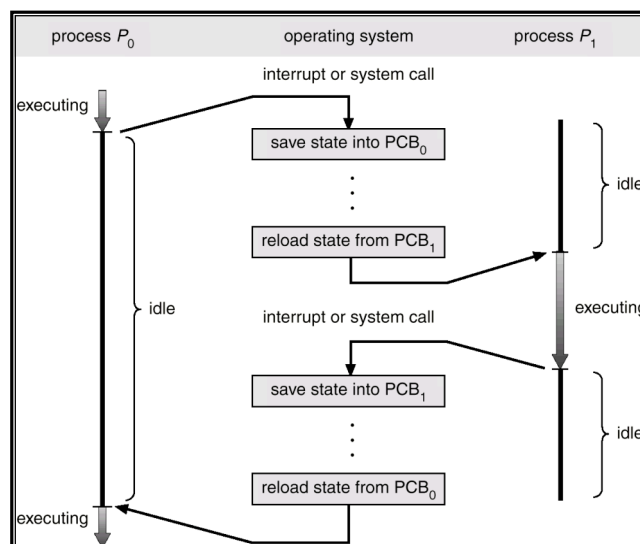
## OS Process Control Block

- Stores all of the information about a process
- PCB contains
  - Process state: new, ready, etc.
  - (saved) processor registers
  - Memory Management Information
    - page tables, limit registers for segments
  - CPU scheduling information
    - process priority; pointers to process queues
  - Accounting information
    - time used (and limits); files used; program owner; parent process id
  - I/O status information
    - list of open files; pending I/O operations

## Sample PCB

pointer	process state
process number	
program counter	
registers	
memory limits	
list of open files	
⋮	

## Context Switch using PCBs





## GeekOS PCB (part I)

- **struct Kernel\_Thread**, contains
  - Process Identifier (PID)
  - Scheduling criteria (priority)
  - Accounting info (CPU clock ticks)
  - Kernel stack pointer
    - Context-switch information stored here; i.e. general-purpose registers, program counter, etc.
- User processes in GeekOS are a special kind of kernel thread

## GeekOS PCB (part II)

- **struct User\_Context**, contains
  - Pointer to process (physical) memory
    - Includes all the segments you set up in Project 1; i.e. code segment, data segment, etc.
  - Pointers to initial program entry point, initial argument, and initial stack.
  - Information for managing address space protection
    - i386 descriptor tables and selectors
- You'll set this up in Project 2

## Storing PCBs


- Track which processes are in which states
  - Collection of PCBs is called a **process table**
- How to store the process table?
- First Option:

P1	P2	P2	P3	P4	P5
Ready	Waiting	New	Term	Waiting	Ready

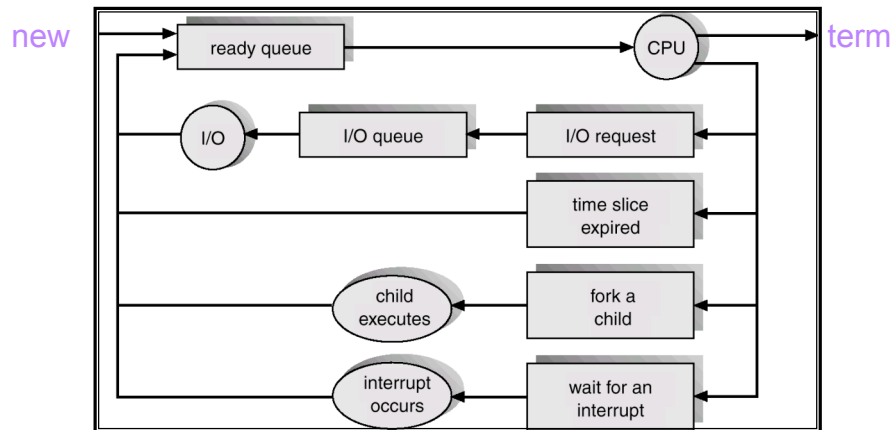
- Simple, but slow to find processes
- Also need additional datastructures for fairness

## Queues of Processes

- Store processes in state-based queues



## State Transitions with Queues



## Forking a New Process

- New process is called the **child**; the process that created it is the **parent**
- Create a PCB for the new process
  - copy most entries from the parent
  - clear accounting fields
  - buffer pending I/O
  - allocate a *PID* for the new process

## Forking a New Process

- Allocate memory for it
  - Might copy all of the parents' segments
  - Text segment could be shared
    - rarely changes
  - Use memory mapping hardware to help
    - will talk more about this in the memory management part of the class
- Add it to the ready queue

## Process Termination

- Process can terminate itself
  - via the *exit* system call
- One process can terminate another process
  - use the *kill* system call
  - can any process kill any other process?
    - No, that would be bad.
    - Normally an ancestor can terminate a descendant
- OS kernel can terminate a process
  - exceeds resource limits
  - tries to perform an illegal operation

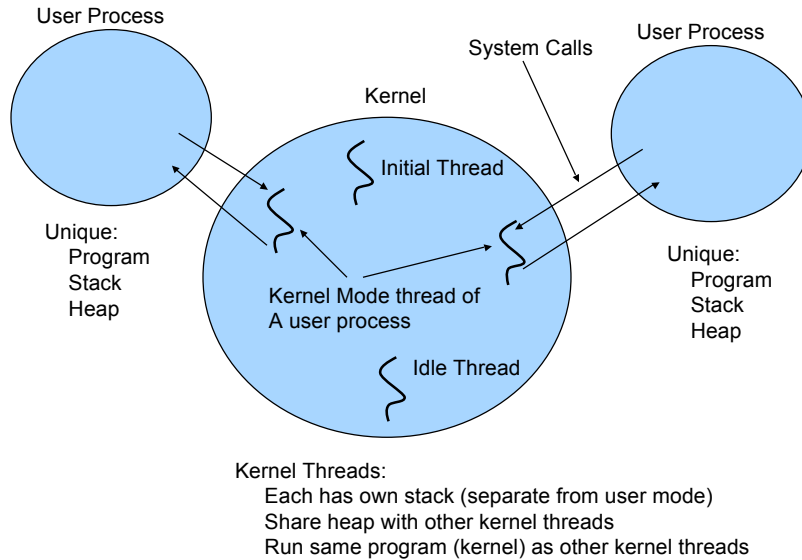
## Orphan Processes

- What if a parent terminates before the child?
  - the child called an **orphan** process
    - in UNIX - becomes child of the root process
    - in VMS - terminated

## UNIX example

- Terminated process
  - signals parent of its death (SIGCHLD)
  - is called a zombie in UNIX
  - remains around waiting to be reclaimed
- Parent process
  - wait system call retrieves info about the dead process
    - exit status
    - accounting information
  - signal handler is generally called the *reaper*
    - since its job is to collect the dead processes


# Kernel Mode and User Mode



## Threads

- Processes can be a heavy (expensive) object
- Threads are like processes but generally a collection of threads will share
  - memory (except stack)
  - open files (and buffered data)
  - signals
- Can be user or system level
  - user level: kernel sees one process
    - + easy to implement by users
    - I/O management is difficult
    - in an multi-processor can't get parallelism
  - system level: kernel schedules threads

# Single and Multithreaded Processes




## Execution Abstractions

- Kernel Threads
  - Threads that run with kernel privileges
- User Threads
  - Threads running in user space
  - Kernel may not be aware of them
- Processes
  - An execution context *with an address space*
  - Visible to and scheduled by the kernel
- Light-Weight Processes
  - An execution context sharing an address space
  - Visible to and scheduled by the kernel

## Multithreading Models


- Many-to-One
- One-to-One
- Many-to-Many

### Many-to-One Model






## One-to-one Model



## Many-to-Many Model



## Why multiple threads?

- Performance:
  - Parallelism on multiprocessors
  - Concurrency of computation and I/O
- Can easily express some programming paradigms
  - Event processing
  - Simulations
- Keep computations separate

## Why not multiple threads?

- Complexity:
  - Dealing with safety, liveness, composition
- Overhead
  - Higher resource usage
- Check out CMSC 433 for lots of information on threads and their alternatives!