

# CMSC 412

Memory Management:  
Paging and Virtual Memory

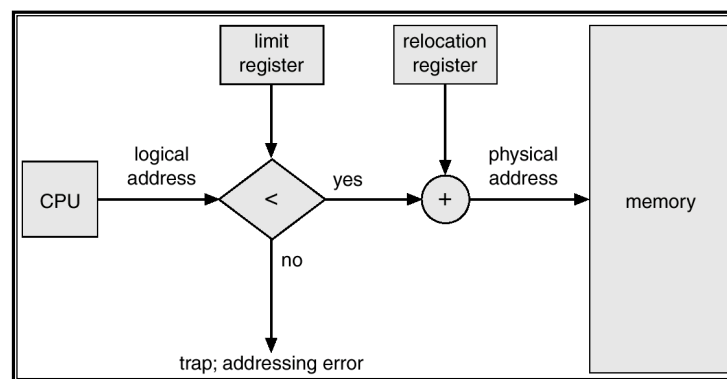
## Announcements

- Reading:
  - Today: Chapter 9.2-9.4
  - Next time: Chapter 9.5-9.6, 10

## Logical vs. Physical Addresses

- The concept of a logical *address space* that is bound to a separate *physical address space* is central to proper memory management.
  - *Logical address* - generated by the CPU; also referred to as *virtual address*.
  - *Physical address* - address seen by the memory unit.
- Just as in Project 2.


## Basic Relocation Hardware



# Swapping

- Main memory is big, but what if we run out?
- A process can be *swapped* temporarily out of memory to a *backing store*, and then brought back into memory for continued execution.
  - Backing store (disk) *bigger* than main memory
- Major part of swap time is transfer time; total transfer time is directly proportional to the *amount* of memory swapped.
  - But disk is *slower* than main memory
- Swapping used on UNIX, Linux, Windows, ...

## Schematic View of Swapping




## Contiguous Allocation

- Main memory usually into two partitions:
  - Resident operating system, usually held in low memory with interrupt vector.
  - User processes then held in high memory.
- Single-partition allocation
  - Subdivide user-space into fixed-size chunks of memory. Each process gets a full chunk.
    - What if the process doesn't need that much?
    - What if it needs more?

## Contiguous Allocation

- Multiple-partition allocation
  - *Hole* - block of available memory; holes of various size are scattered throughout memory.
  - When a process arrives, it is allocated memory from a hole large enough to accommodate it.
  - Operating system maintains information about:
    - a) allocated partitions    b) free partitions (hole)



## Dynamic Storage-Allocation Problem

- How to satisfy a request of size  $n$  from a list of free holes?
- Goal: efficient utilization of storage.
  - If we have  $n$  processes that require  $m$  bytes of memory where  $m < n$ , we can fit all processes in memory at once.
- Goal: fast allocation times.
  - The time to find a hole of necessary size

## Dynamic Storage-Allocation Problem

- **First-fit**: Allocate the first hole that is big enough.
- **Best-fit**: Allocate the smallest hole that is big enough; must search entire list, unless ordered by size. Produces the smallest leftover hole.
- **Worst-fit**: Allocate the largest hole; must also search entire list. Produces the largest leftover hole.
- First-fit and best-fit better than worst-fit in terms of speed and storage utilization.

## Fragmentation

- **External Fragmentation** - total memory space exists to satisfy a request, but it is not contiguous.
- **Internal Fragmentation** - allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used.

## Reducing Fragmentation

- Compaction
  - Shuffle memory contents to place all free memory together in one large block.
- Compaction is possible *only* if using relocatable logical addresses.
- I/O problem
  - Pin job in memory while it is doing I/O.
  - Or, do I/O only into OS buffers.

## Non-Contiguous Allocation

- Can eliminate external fragmentation and improve storage utilization by allowing process memory to be *non-contiguous*.
  - Break the process memory into different chunks and allocate chunks in different parts of physical memory.
- Could you do this using the scheme you have implemented for GeekOS so far?

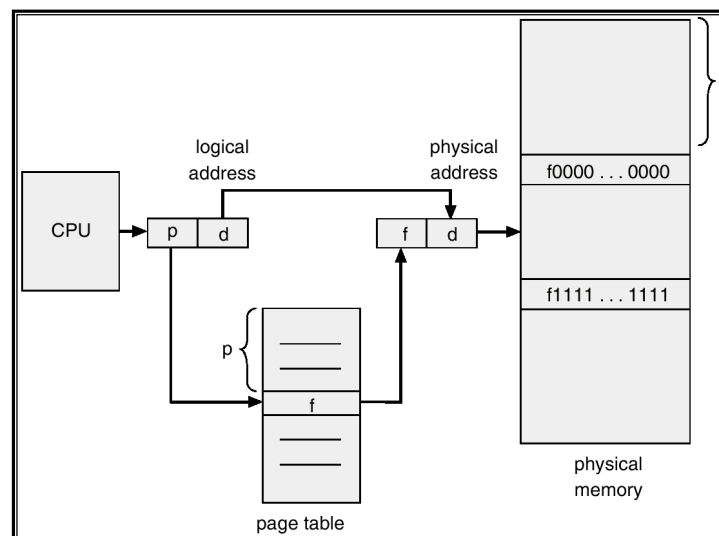
## Paging

- Divide physical memory into fixed sized chunks called **frames** (pages), logical memory into same-sized **pages**
  - typical frames are 512 bytes to 64 Kbytes
  - When a process is to be executed, load the pages that are needed into memory.
- Have a map from pages to frames
  - Called the **page table**.

# Address Translation Scheme

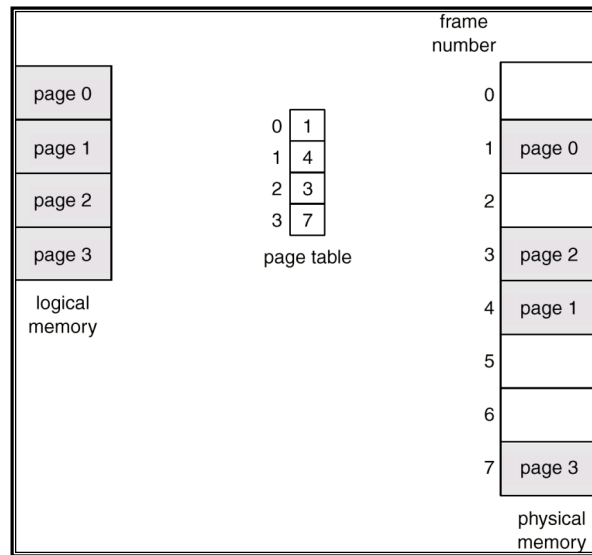
- Address generated by CPU is divided into:
  - *Page number (p)* - used as an index into the page table, which contains base address of each page in physical memory.
  - *Page offset (d)* - combined with base address to define the physical memory address that is sent to the memory unit.
- Consider 32-bit addresses on Pentium:
  - 4096-byte pages (12 bits for the offset)
  - 20 bits for the page number

## Address Translation






# Paging Example



# Free Frames



## Paging Fragmentation

- If we have 4K pages, and a process requires 15K of memory, how many pages will we need?
- Suffer from internal fragmentation: last page likely not completely used
  - On average 50% of a page lost per process to internal fragmentation
- How does this impact our choice of page size?

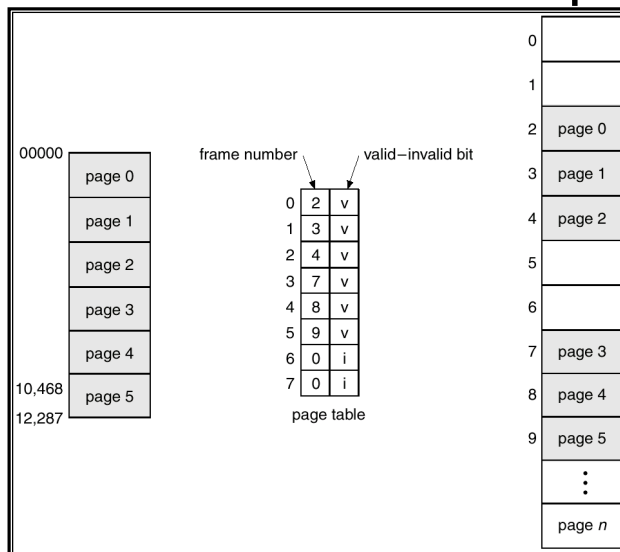
## Implementation of Page Table

- Page table is kept in main memory.
- *Page-table base register* (PTBR) points to the page table.

# Memory Protection

- Memory protection by associating **protection bit** with each frame.
- **Valid-invalid bit** attached to each entry in the page table:
  - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page.
  - “invalid” indicates that the page is not in the process’ logical address space.

## Valid/Invalid Bit Example



## Problem: Page Table Size

- One page table can get very big
  - $2^{20}$  entries (for most programs, most items are empty)
- Simple Solution
  - *Page-table length register* (PRLR) indicates size of the page table.
- But not as good for “sparse” address spaces

## Hierarchical Page Tables

- Break up the logical address space into multiple page tables.
- Simplest case: a two-level page table.
  - Used on the Pentium. The outer table is called the [page directory](#), which points to individual [page tables](#). Each page table is itself a page.

## Two-Level Paging Example

- A logical address (on 32-bit machine with 4K page size) is divided into:
  - a page number consisting of 20 bits.
  - a page offset consisting of 12 bits.
- Since the page table is paged, the page number is further divided into:
  - a 10-bit page number.
  - a 10-bit page offset.

## Two-Level Paging Example

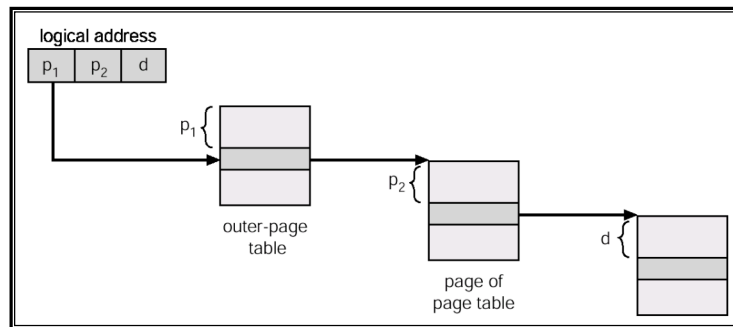
- Thus, a logical address is as follows:

page number		page offset
$p_1$	$p_2$	$d$
10	10	12

where  $p_1$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the outer page table.

# Address-Translation Scheme

- Address-translation scheme for a two-level 32-bit paging architecture



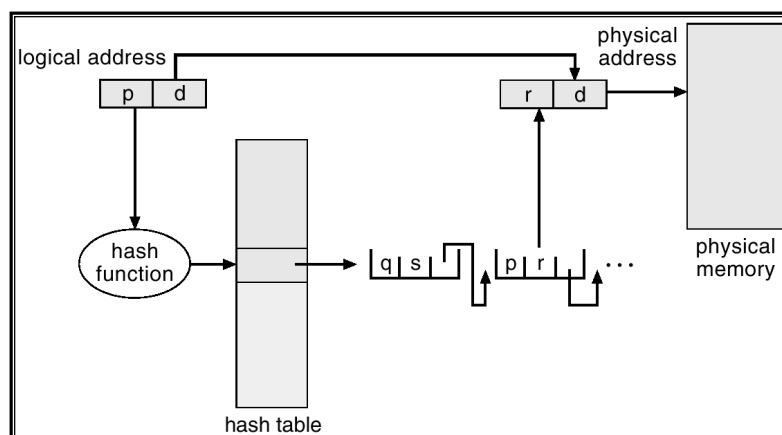
## Problem: Still too big?

- What about 64 bit address spaces?
- Two-level page table:
  - Page num: 52 bits, or  $2^{52}$  entries in page directory!
- More levels of hierarchy?
  - Further subdivide the page number.
    - 32-bit SPARC: three-level scheme.
    - 32-bit 68030: four-level scheme
- Try something different ...

# Hashed Page Tables

- Common in address spaces > 32 bits.
- The virtual page number is hashed into a page table. This page table contains a chain of elements hashing to the same location.
- Virtual page numbers are compared in this chain searching for a match. If a match is found, the corresponding physical frame is extracted.

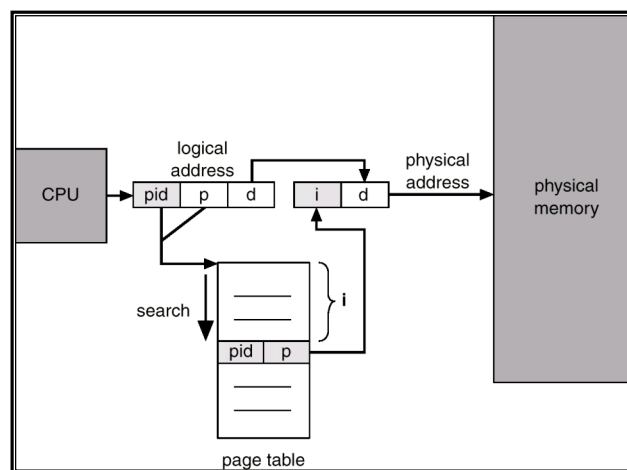
## Hashed Page Table



## Inverted Page Table

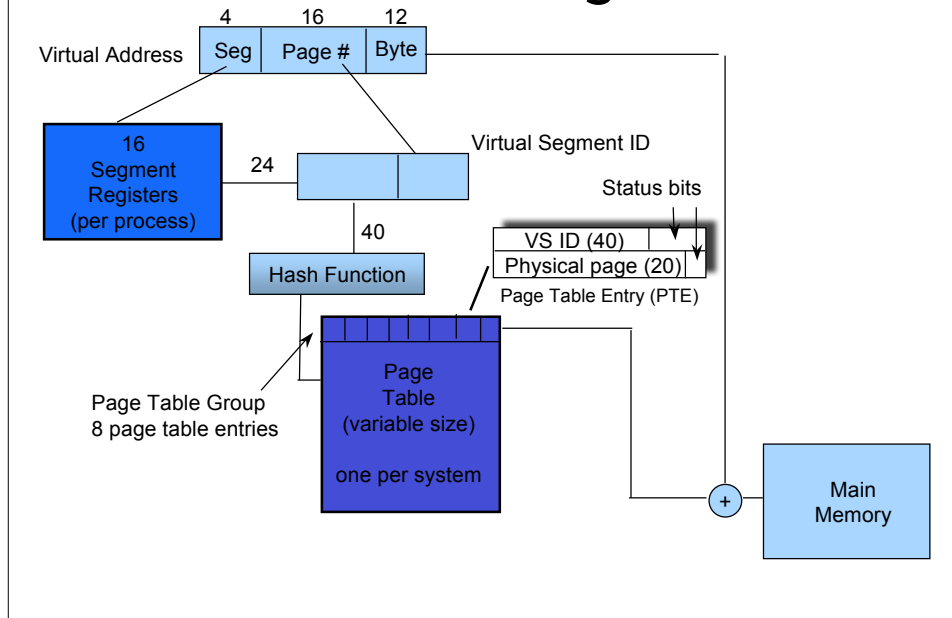
- One entry for each real page of memory.
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page.
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs.
- Use hash table to limit the search to one — or at most a few — page-table entries.

## Inverted Page Table Architecture






## PPC Inverted Page Table



## Problem: Address Lookup

- Every data/instruction access requires many memory accesses. At the least, one for the page table and one for the data/instruction.
- The *two memory access problem* can be solved by the use of a special fast-lookup hardware cache called *associative memory* or *translation look-aside buffers (TLBs)*
  - typically 16- 64 entries

## Paging Hardware With TLB



## Example: i386 TLB

Valid	Virtual Page	Physical Page

20 bits      20 bits

## Effective Access Time

- Associative Lookup =  $\epsilon$  time unit
- Assume memory cycle time is 1 microsecond
- Hit ratio - percentage of times that a page number is found in the associative registers; ration related to number of associative registers.
- Hit ratio =  $\alpha$
- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} &= (1 + \epsilon) \alpha + (2 + \epsilon)(1 - \alpha) \\ &= 2 + \epsilon - \alpha \end{aligned}$$

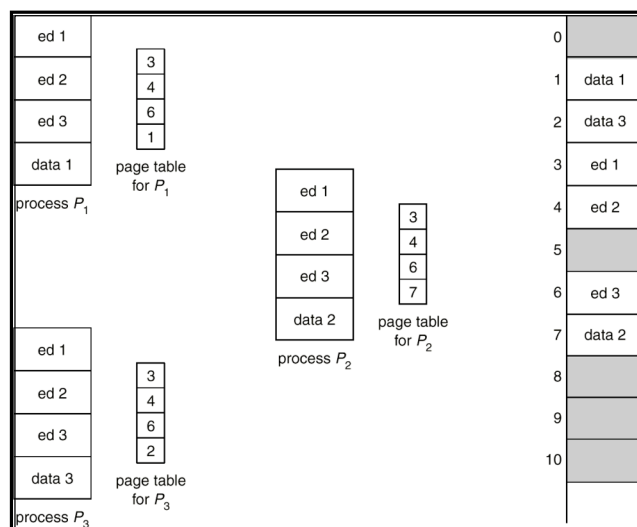
## Super Pages

- TLB Entries
  - Tend to be limited in number
  - Can only refer to one page
- Idea
  - Create bigger pages
  - 4MB instead of 4KB
  - One TLB entry covers more memory
  - What's the tradeoff?

# Shared Pages

- Shared code
  - One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems).
  - Shared code must appear in same location in the logical address space of all processes.
- Private code and data
  - Each process keeps a separate copy of the code and data.
  - The pages for the private code and data can appear anywhere in the logical address space.

## Shared Pages Example



## Sharing Writable Pages

- When writes occur, decide if processes share data
  - operating systems often implement *copy on write* - pages are shared until a process carries out a write
    - when a shared page is written, a new page frame is allocated
    - writing process owns the modified page
    - all other sharing processes own the original
  - processes use semaphores or other means to coordinate access

## Virtual Memory

- Paging is a typical feature of broader support for *virtual memory*
- Pages can be swapped to and from disk
  - Presents the illusion of a larger physical memory
- Virtual addressing simplifies model
  - Swapping: can load pages into different addresses in physical memory
  - Addressing: different processes have their own address spaces, starting at 0, separate from other processes